

# Structured Query Language

***MSBA 7024 / MACC 7020***

***Database Design and Management***

# Objectives

- Definition of terms
- Interpret history and role of SQL
- Define a database using SQL data definition language
- Write single table queries using SQL
- Establish referential integrity using SQL
- Write multiple table SQL queries
- Define and use different types of joins
- Write noncorrelated and correlated subqueries
- Establish transaction integrity in SQL
- Understand triggers and stored procedures

# SQL Overview

- Structured Query Language
- The standard for relational database management systems (RDBMS)
- RDBMS: A database management system that manages data as a collection of tables in which all relationships are represented by common values in related tables

# History of SQL

- 1970: E. Codd develops relational database concept
- 1974-1979: System R with Sequel (later SQL) created at IBM Research Lab
- 1979: Oracle markets first relational DB with SQL
- 1986: ANSI SQL standard released
- 1989, 1992, 1999, 2003, 2006, 2008, 2011, 2016: Major ANSI standard updates
- Current: SQL is supported by all major relational database vendors

# Purpose of SQL Standard

- Specify syntax/semantics for data definition and manipulation
- Define data structures
- Enable portability
- Specify minimal (level 1) and complete (level 2) standards
- Allow for later growth/enhancement to standard

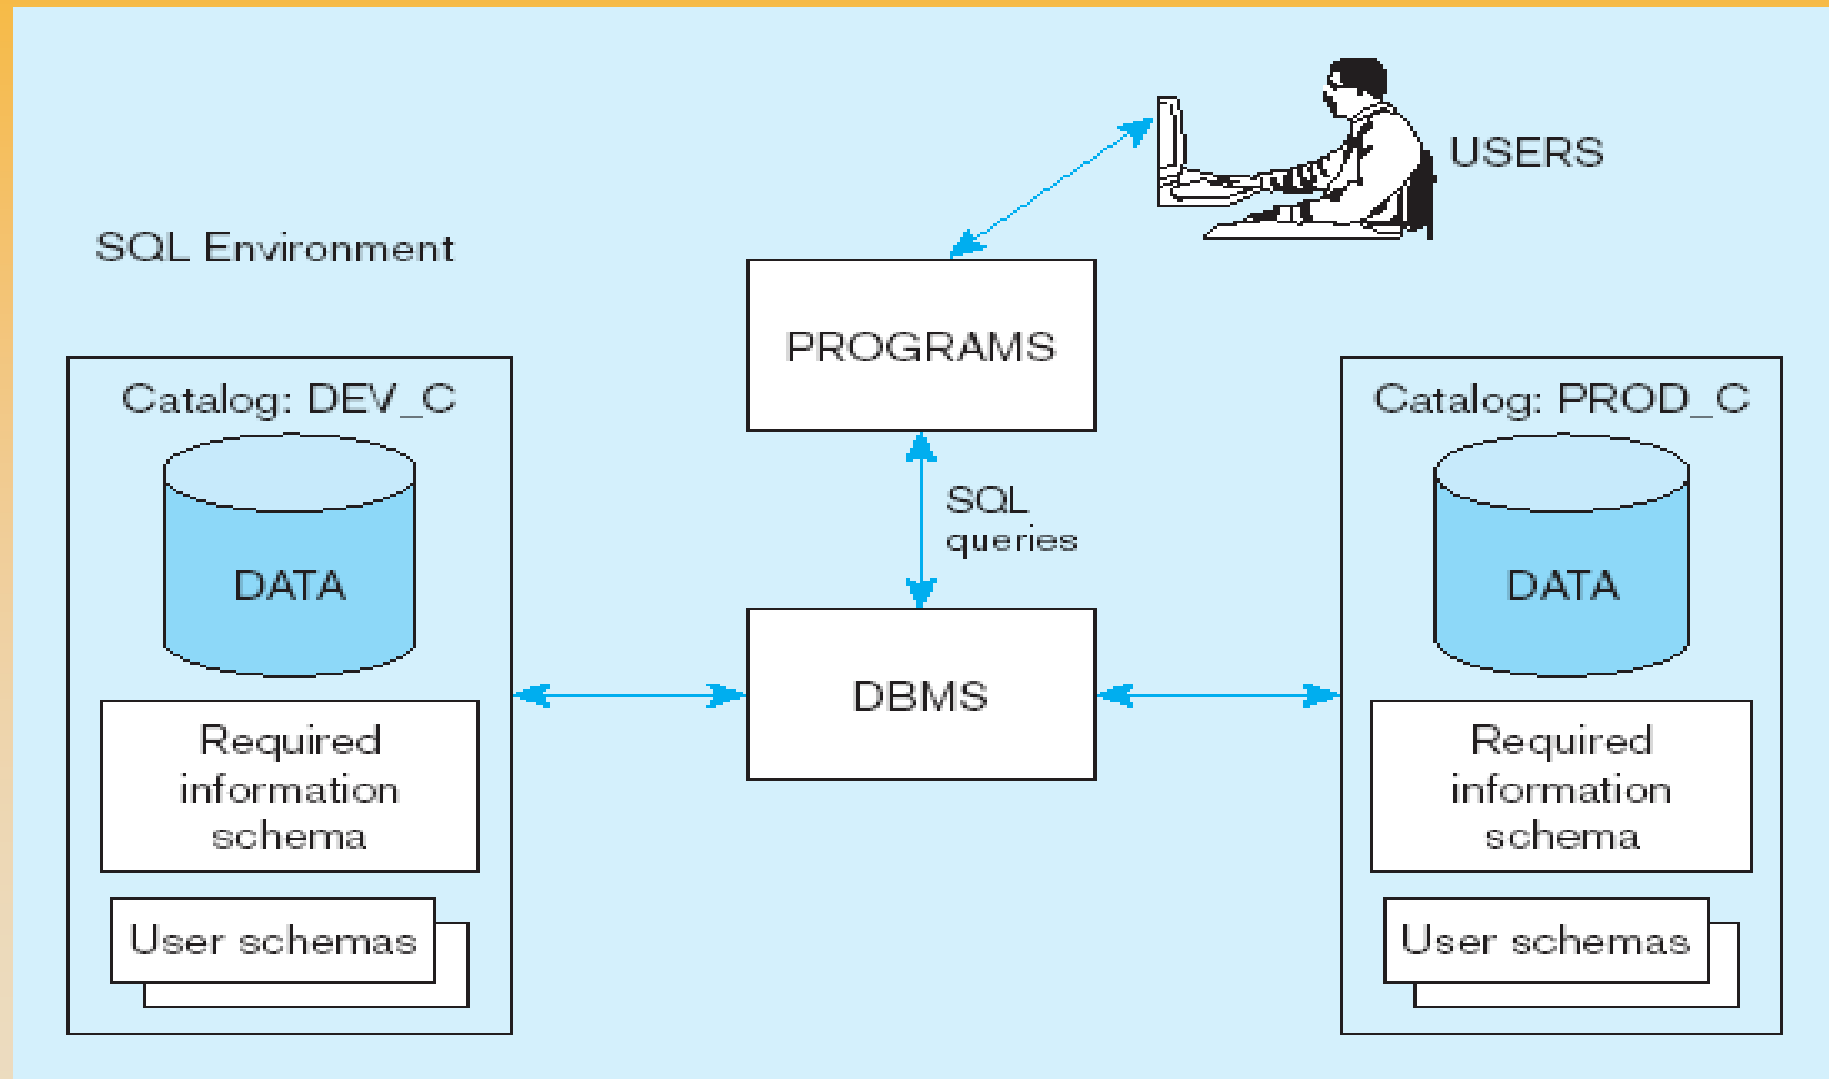
# Benefits of a Standardized Relational Language

- Reduced training costs
- Productivity
- Application portability
- Application longevity
- Reduced dependence on a single vendor
- Cross-system communication

# SQL Environment

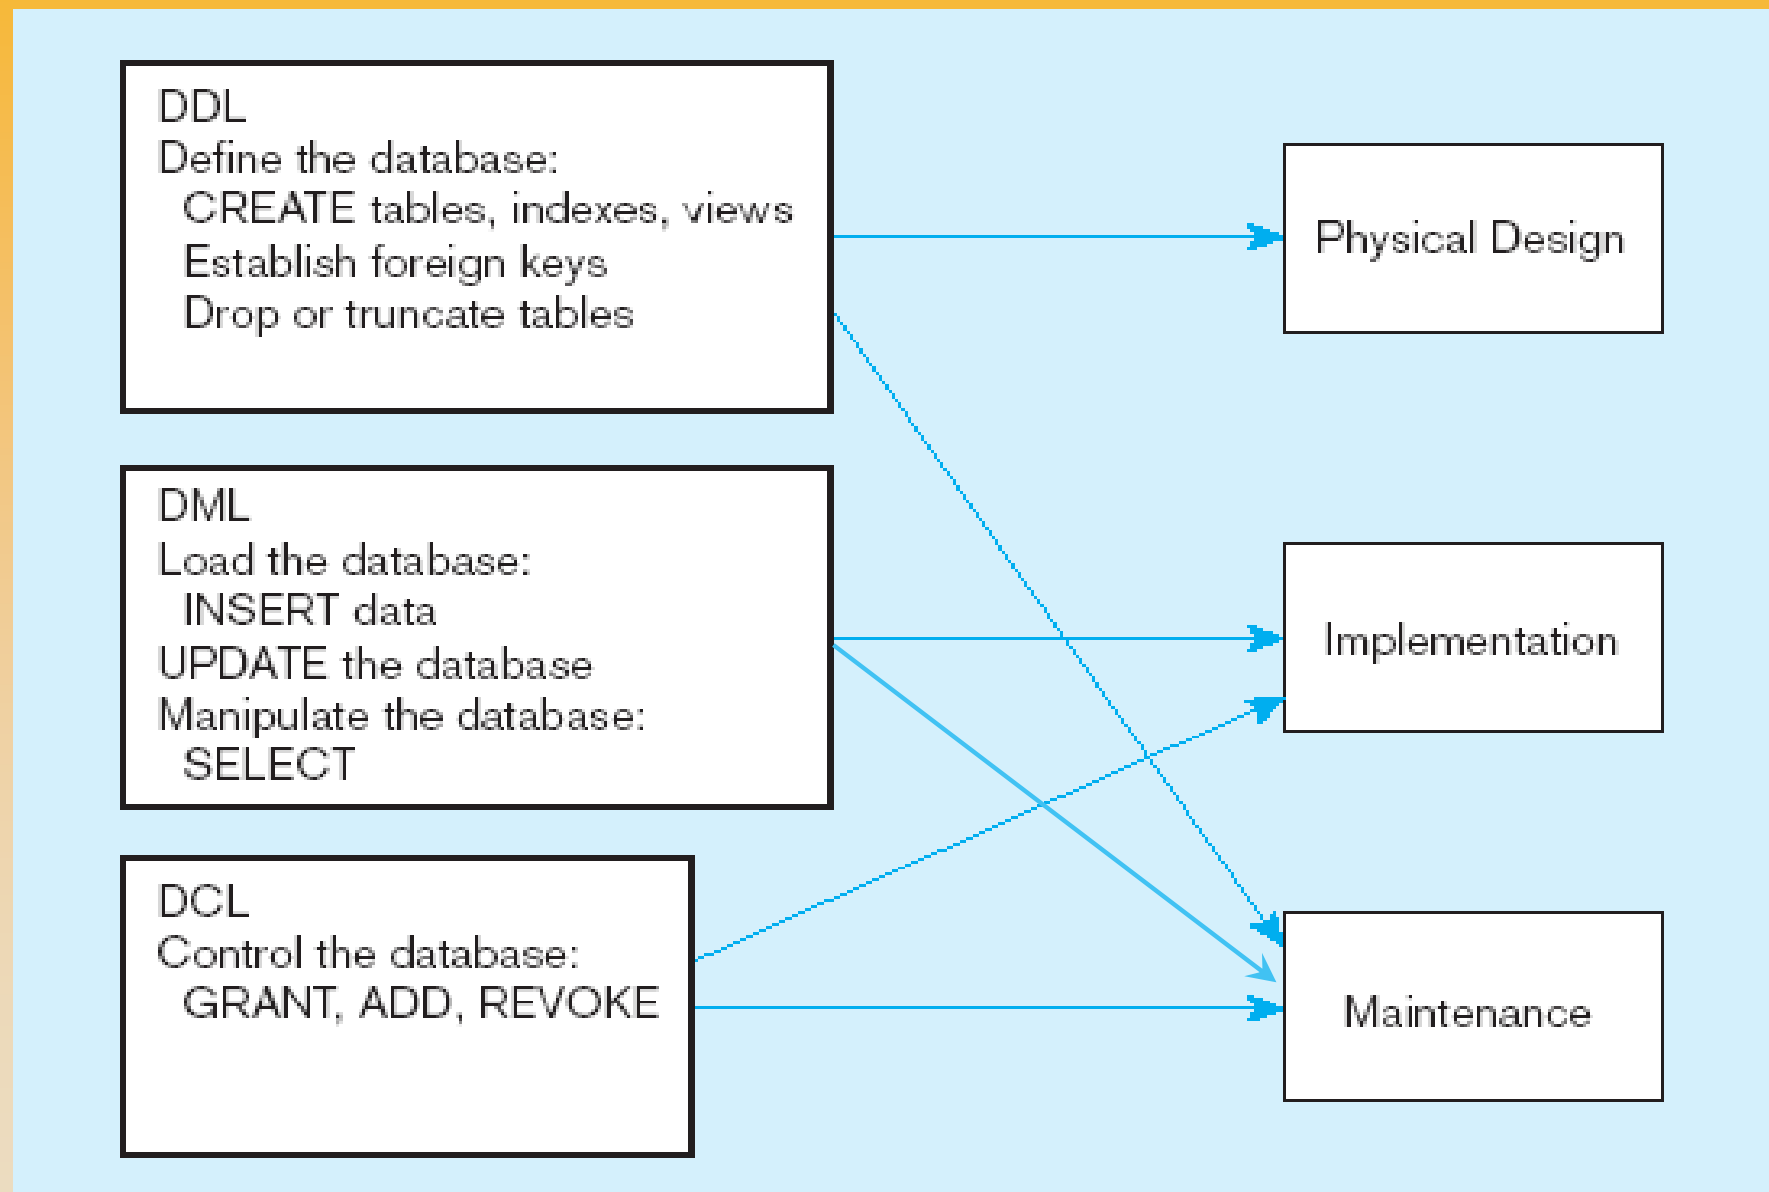
- Catalog
  - A set of schemas that constitute the description of a database
- Schema
  - The structure that contains descriptions of objects created by a user (base tables, views, constraints)
- Data Definition Language (DDL)
  - Commands that define a database, including creating, altering, and dropping tables and establishing constraints
- Data Manipulation Language (DML)
  - Commands that maintain and query a database
- Data Control Language (DCL)
  - Commands that control a database, including administering privileges and committing data

# A simplified schematic of a typical SQL environment





# DDL, DML, DCL, and the database development process



# SQL Database Definition

- Data Definition Language (DDL)
- Major CREATE statements:
  - CREATE SCHEMA—defines a portion of the database owned by a particular user
  - CREATE TABLE—defines a table and its columns
  - CREATE VIEW—defines a logical table from one or more views
- Other CREATE statements: CHARACTER SET, COLLATION, TRANSLATION, ASSERTION, DOMAIN

# Table Creation

## General syntax for CREATE TABLE

```
CREATE TABLE tablename
( {column definition [table constraint] } . . .
[ON COMMIT {DELETE | PRESERVE} ROWS] );

where column definition ::=
column_name
    {domain name | datatype [(size)] }
    [column_constraint_clause . . .]
    [default value]
    [collate clause]

and table constraint ::=
    [CONSTRAINT constraint_name]
    Constraint_type [constraint_attributes]
```

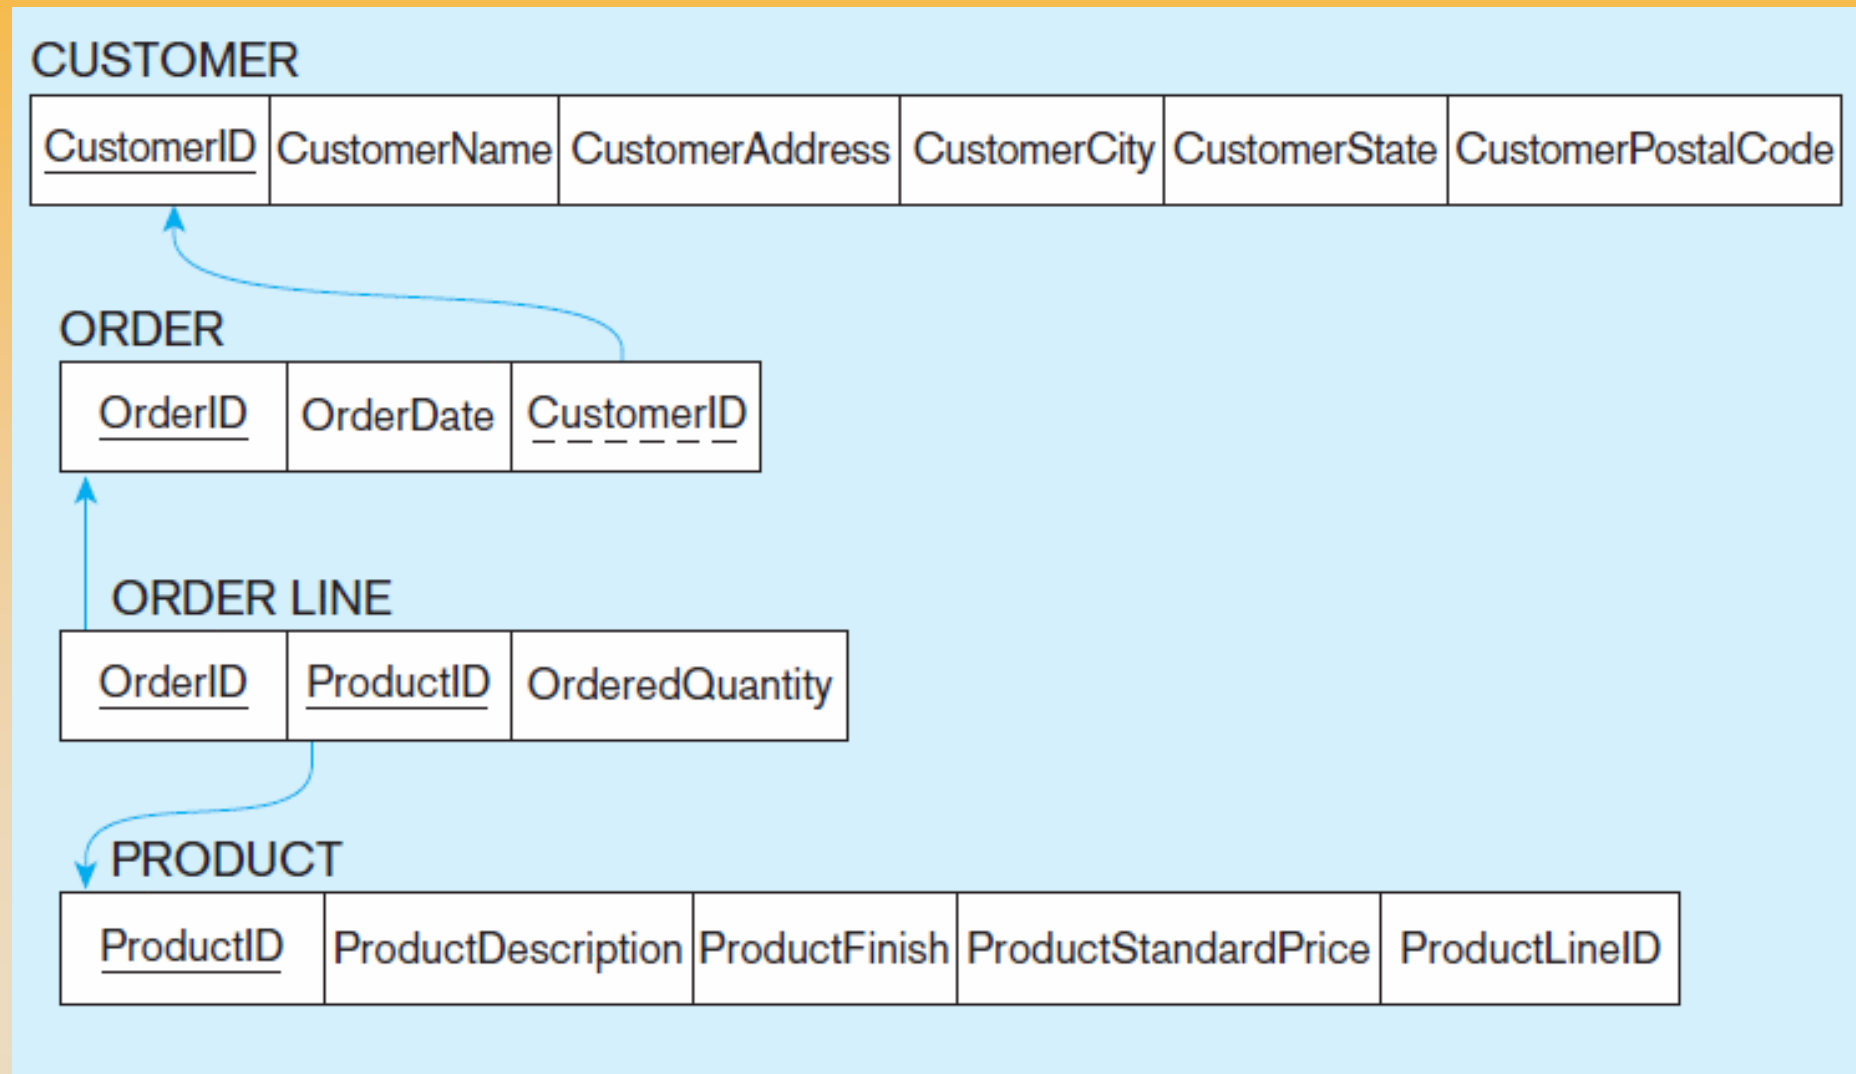
## Steps in table creation:

1. Identify data types for attributes
2. Identify columns that can and cannot be null
3. Identify columns that must be unique (candidate keys)
4. Identify primary key–foreign key mates
5. Determine default values
6. Identify constraints on columns (domain specifications)
7. Create the table and associated indexes

# Some SQL Data types

String	CHARACTER (CHAR)	Stores string values containing any characters in a character set. CHAR is defined to be a fixed length.
	CHARACTER VARYING (VARCHAR)	Stores string values containing any characters in a character set, but of definable variable length.
	BINARY LARGE OBJECT (BLOB)	Stores binary string values in hexadecimal format. BLOB is defined to be a variable length.
Number	NUMERIC	Stores exact numbers with a defined precision and scale.
	INTEGER (INT)	Stores exact numbers with a predefined precision and scale of zero.
Temporal	TIMESTAMP	Stores a moment an event occurs, using a definable fraction of a second precision.
Boolean	BOOLEAN	Stores truth values, TRUE, FALSE, or UNKNOWN.

# The following slides create tables for these table designs



# SQL database definition commands for Pine Valley Furniture

## Overall table definitions

```
CREATE TABLE Customer_T
    (CustomerID          NUMBER(11,0)    NOT NULL,
     CustomerName        VARCHAR2(25)    NOT NULL,
     CustomerAddress     VARCHAR2(30),
     CustomerCity        VARCHAR2(20),
     CustomerState       CHAR(2),
     CustomerPostalCode  VARCHAR2(9),
 CONSTRAINT Customer_PK PRIMARY KEY (CustomerID));
```

```
CREATE TABLE Order_T
    (OrderID             NUMBER(11,0)    NOT NULL,
     OrderDate           DATE DEFAULT SYSDATE,
     CustomerID          NUMBER(11,0),
 CONSTRAINT Order_PK PRIMARY KEY (OrderID),
 CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID));
```

```
CREATE TABLE Product_T
    (ProductID           NUMBER(11,0)    NOT NULL,
     ProductDescription   VARCHAR2(50),
     ProductFinish       VARCHAR2(20)
                          CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                                                    'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice DECIMAL(6,2),
     ProductLineID       INTEGER,
 CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

```
CREATE TABLE OrderLine_T
    (OrderID             NUMBER(11,0)    NOT NULL,
     ProductID           INTEGER         NOT NULL,
     OrderedQuantity     NUMBER(11,0),
 CONSTRAINT OrderLine_PK PRIMARY KEY (OrderID, ProductID),
 CONSTRAINT OrderLine_FK1 FOREIGN KEY (OrderID) REFERENCES Order_T(OrderID),
 CONSTRAINT OrderLine_FK2 FOREIGN KEY (ProductID) REFERENCES Product_T(ProductID));
```

# Defining attributes and their data types

```
CREATE TABLE Product_T
```

(ProductID	NUMBER(11,0)	NOT NULL,
ProductDescription	VARCHAR2(50),	
ProductFinish	VARCHAR2(20)	

```
        CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',  
                                'Red Oak', 'Natural Oak', 'Walnut')),
```

ProductStandardPrice	DECIMAL(6,2),
ProductLineID	INTEGER,

```
CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

## Non-nullable specification

```
CREATE TABLE Product_T
    (ProductID                NUMBER(11,0)    NOT NULL,
     ProductDescription        VARCHAR2(50),
     ProductFinish             VARCHAR2(20)
                                CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                                                         'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice     DECIMAL(6,2),
     ProductLineID             INTEGER,
    CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

Primary keys  
can never have  
NULL values

## Identifying primary key



## Non-nullable specifications

```
CREATE TABLE OrderLine_T
```

```
    (OrderID
```

```
        NUMBER(11,0)
```

```
    NOT NULL,
```

```
    ProductID
```

```
        INTEGER
```

```
    NOT NULL,
```

```
    OrderedQuantity
```

```
        NUMBER(11,0),
```

```
    CONSTRAINT OrderLine_PK PRIMARY KEY (OrderID, ProductID),
```

**Primary key**

```
    CONSTRAINT OrderLine_FK1 FOREIGN KEY (OrderID) REFERENCES Order_T(OrderID),
```

```
    CONSTRAINT OrderLine_FK2 FOREIGN KEY (ProductID) REFERENCES Product_T(ProductID));
```

Some primary keys are composite—  
composed of multiple attributes

# Controlling the values in attributes

```
CREATE TABLE Order_T
    (OrderID                NUMBER(11,0)    NOT NULL,
     OrderDate              DATE DEFAULT SYSDATE,
     CustomerID             NUMBER(11,0),
 CONSTRAINT Order_PK PRIMARY KEY (OrderID),
 CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID));

CREATE TABLE Product_T
    (ProductID              NUMBER(11,0)    NOT NULL,
     ProductDescription      VARCHAR2(50),
     ProductFinish           VARCHAR2(20)
     CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                              'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice   DECIMAL(6,2),
     ProductLineID          INTEGER,
 CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

**Default value**

**Domain constraint**

# Identifying foreign keys and establishing relationships

```
CREATE TABLE Customer_T
```

(CustomerID	NUMBER(11,0)	NOT NULL,
CustomerName	VARCHAR2(25)	NOT NULL,
CustomerAddress	VARCHAR2(30),	
CustomerCity	VARCHAR2(20),	
CustomerState	CHAR(2),	
CustomerPostalCode	VARCHAR2(9),	

Primary key of  
parent table

```
CONSTRAINT Customer_PK PRIMARY KEY (CustomerID));
```

```
CREATE TABLE Order_T
```

(OrderID	NUMBER(11,0)	NOT NULL,
OrderDate	DATE DEFAULT SYSDATE,	
CustomerID	NUMBER(11,0),	

```
CONSTRAINT Order_PK PRIMARY KEY (OrderID),
```


```
CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID));
```

Foreign key of dependent table

# Data Integrity Controls

- Referential integrity—constraint that ensures that foreign key values of a table must match primary key values of a related table in 1:M relationships
- Restricting:
  - Deletes of primary records
  - Updates of primary records
  - Inserts of dependent records

# Ensuring data integrity through updates



**Restricted Update:** A customer ID can only be deleted if it is not found in ORDER table.

```
CREATE TABLE CustomerT
    (CustomerID          INTEGER DEFAULT '999'    NOT NULL,
     CustomerName        VARCHAR(40)             NOT NULL,
     ...
    CONSTRAINT Customer_PK PRIMARY KEY (CustomerID),
    ON UPDATE RESTRICT);
```

**Cascaded Update:** Changing a customer ID in the CUSTOMER table will result in that value changing in the ORDER table to match.

```
... ON UPDATE CASCADE);
```

**Set Null Update:** When a customer ID is changed, any customer ID in the ORDER table that matches the old customer ID is set to NULL.

```
... ON UPDATE SET NULL);
```

**Set Default Update:** When a customer ID is changed, any customer ID in the ORDER tables that matches the old customer ID is set to a predefined default value.

```
... ON UPDATE SET DEFAULT);
```

Relational  
integrity is  
enforced via  
the primary-  
key to foreign-  
key match

# Changing and Removing Tables

- **ALTER TABLE** statement allows you to change column specifications:

Syntax:

- **ALTER TABLE** table\_name alter\_table\_action;
  - **ADD [COLUMN]** column\_definition
  - **ALTER [COLUMN]** column\_name **SET DEFAULT** default-value
  - **ALTER [COLUMN]** column\_name **DROP DEFAULT**
  - **DROP [COLUMN]** column\_name [**RESTRICT**] [**CASCADE**]
  - **ADD** table\_constraint

Command: To add a customer type column named CustomerType to the CUSTOMER table, set default value as "Commercial".

- **ALTER TABLE** Customer\_T  
    **ADD COLUMN** CustomerType VARCHAR(12) **DEFAULT** "Commercial";

# Changing and Removing Tables

- DROP TABLE statement allows you to remove tables from your schema:
  - DROP TABLE Customer\_T

# Schema Definition

- Control processing/storage efficiency:
  - Choice of indexes
  - File organizations for base tables
  - File organizations for indexes
  - Data clustering
  - Statistics maintenance
- Creating indexes
  - Speed up random/sequential access to base table data
  - Example
    - `CREATE INDEX CustomerNameIdx ON Customer_T(CustomerName)`
    - This makes an index for the CustomerName field of the Customer\_T table



# Insert Statement

- Adds data to a table
- Inserting into a table
  - `INSERT INTO Customer_T VALUES (001, 'Contemporary Casuals', '1355 S. Himes Blvd.', 'Gainesville', 'FL', 32601);`
- Inserting a record that has some null attributes requires identifying the fields that actually get data
  - `INSERT INTO Product_T (ProductID, ProductDescription, ProductFinish, ProductStandardPrice) VALUES (1, 'End Table', 'Cherry', 175);`
- Inserting from another table
  - `INSERT INTO CA_Customer_T SELECT * FROM Customer_T WHERE CustomerState = 'CA';`

# Creating Tables with Identity Columns

```
CREATE TABLE Customer_T  
(CustomerID INTEGER GENERATED ALWAYS AS IDENTITY  
  (START WITH 1  
   INCREMENT BY 1  
   MINVALUE 1  
   MAXVALUE 10000  
   NO CYCLE),  
 CustomerName          VARCHAR2(25) NOT NULL,  
 CustomerAddress       VARCHAR2(30),  
 CustomerCity          VARCHAR2(20),  
 CustomerState         CHAR(2),  
 CustomerPostalCode    VARCHAR2(9),  
 CONSTRAINT Customer_PK PRIMARY KEY (CustomerID);
```

Introduced with SQL:200n

Inserting into a table does not require explicit customer ID entry or field list

```
INSERT INTO Customer_T VALUES ('Contemporary Casuals', '1355  
S. Himes Blvd.', 'Gainesville', 'FL', 32601);
```

# Delete Statement

- Removes rows from a table
- Delete certain rows
  - `DELETE FROM Customer_T WHERE CustomerState = 'HI';`
- Delete all rows
  - `DELETE FROM Customer_T;`

# Update Statement

- Modifies data in existing rows
- ```
UPDATE Product_T  
SET ProductStandardPrice = 775  
WHERE ProductID= 7;
```

# SELECT Statement

- Used for queries on single or multiple tables
- Clauses of the SELECT statement:
  - **SELECT**
    - List the columns (and expressions) that should be returned from the query
  - **FROM**
    - Indicate the table(s) or view(s) from which data will be obtained
  - **WHERE**
    - Indicate the conditions under which a row will be included in the result
  - **GROUP BY**
    - Indicate categorization of results
  - **HAVING**
    - Indicate the conditions under which a category (group) will be included
  - **ORDER BY**
    - Sorts the result according to specified criteria

# SELECT Example

- Find products with standard price less than \$275

```
SELECT ProductID, ProductStandardPrice  
FROM Product_T  
WHERE ProductStandardPrice < 275;
```

Comparison Operators in SQL

| Operator | Meaning                  |
|----------|--------------------------|
| =        | Equal to                 |
| >        | Greater than             |
| >=       | Greater than or equal to |
| <        | Less than                |
| <=       | Less than or equal to    |
| <>       | Not equal to             |
| !=       | Not equal to             |

# SELECT Example

## Using a Function

- Using the COUNT *aggregate function* to find totals

```
SELECT COUNT(*) FROM OrderLine_T  
WHERE OrderID = 1004;
```

Note: with aggregate functions you can't have single-valued columns included in the SELECT clause

```
SELECT COUNT(*), ProductID FROM OrderLine_T  
WHERE OrderID = 1004;
```

The above statement will result in error.

# SELECT Example Using a Function

- Other aggregate functions: MAX, MIN, SUM, AVG

```
SELECT MAX(ProductID) FROM OrderLine_T  
WHERE OrderID = 1004;
```

```
SELECT MIN(ProductDescription)  
FROM Product_T;
```




# SELECT Example–Boolean Operators

- **AND, OR, and NOT** Operators for customizing conditions in WHERE clause
- SELECT ProductDescription, ProductFinish, ProductStandardPrice  
FROM Product\_T  
WHERE (ProductDescription **LIKE** '%Desk'  
**OR** ProductDescription **LIKE** '%Table')  
**AND** ProductStandardPrice > 300;
- Use **\*** instead of **%** in MS-Access

Note: the LIKE operator allows you to compare strings using wildcards. For example, the % wildcard in '%Desk' indicates that all strings that have any number of characters preceding the word "Desk" will be allowed

# Venn Diagram from the Previous Query



# SELECT Example –

## Sorting Results with the ORDER BY clause

- Sort the results first by CustomerState, and within a state by CustomerName

```
SELECT CustomerName, CustomerCity, CustomerState  
FROM Customer_T
```

```
WHERE CustomerState IN ('FL', 'TX', 'CA', 'HI')
```

```
ORDER BY CustomerState, CustomerName;
```

- Use **DESC** for sorting in descending order  
e.g., **ORDER BY** CustomerState **DESC**, CustomerName;

Note: the IN operator in this example allows you to include rows whose CustomerState value is either FL, TX, CA, or HI. It is more efficient than separate OR conditions

# SELECT Example–

## Categorizing Results Using the GROUP BY clause

- For use with aggregate functions
  - ***Scalar aggregate***: single value returned from SQL query with aggregate function
  - ***Vector aggregate***: multiple values returned from SQL query with aggregate function (via GROUP BY)

```
SELECT CustomerState, COUNT(CustomerState)
FROM Customer_T
GROUP BY CustomerState;
```

Note: you can use single-value fields with aggregate functions if they are included in the GROUP BY clause

# SELECT Example–

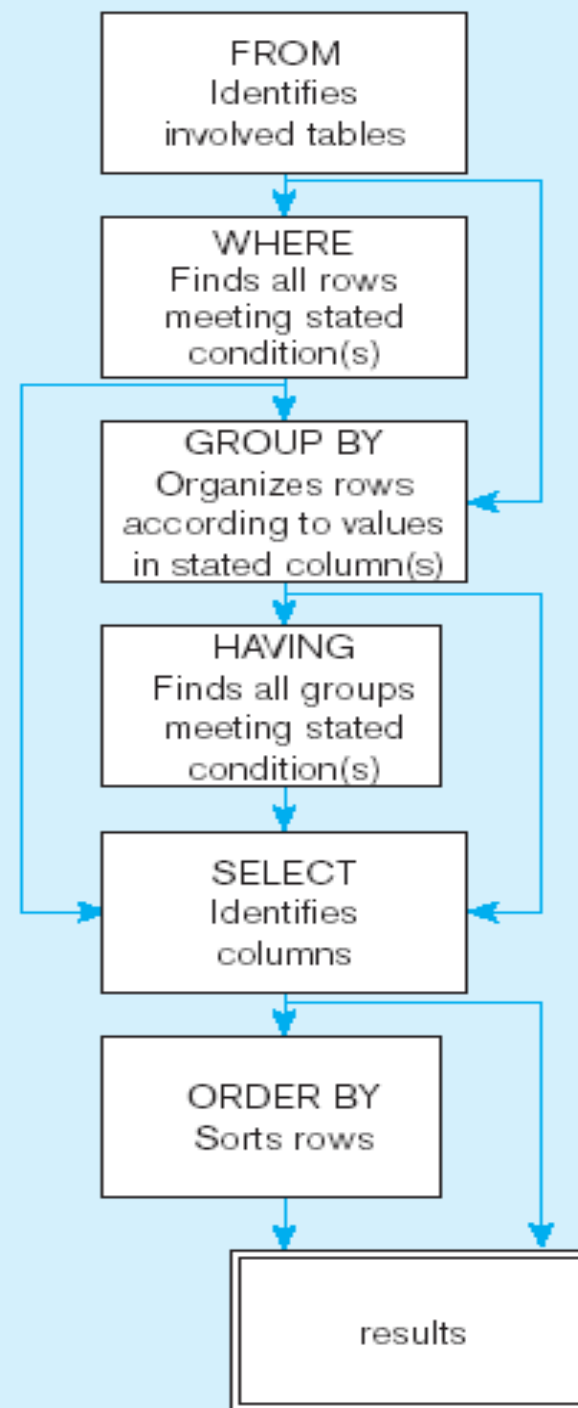
Qualifying Results by Categories  
Using the HAVING Clause

- For use with GROUP BY

```
SELECT CustomerState, COUNT(CustomerState)
FROM Customer_T
GROUP BY CustomerState
HAVING COUNT(CustomerState) > 1;
```

Like a WHERE clause, but it operates on groups (categories), not on individual rows. Here, only those groups with total numbers greater than 1 will be included in final result

## SQL statement processing order (adapted from van der Lans, p.100)



# Using and Defining Views

- Views provide users controlled access to tables
- Base Table – table containing the raw data

```
CREATE VIEW ExpensiveStuff_V
AS
  SELECT ProductID, ProductDescription, ProductStandardPrice
  FROM Product_T
  WHERE ProductStandardPrice > 300
  WITH CHECK OPTION;
```

- View has a name
- View is based on a SELECT statement
- CHECK\_OPTION works only for updateable views and prevents updates that would create rows not included in the view

# Advantages of Views

- Simplify query commands
- Enhance programming productivity
- Assist with data security (but don't rely on views for security, there are more important security measures)
- Provide customized view for user



# Dynamic vs Materialized Views

## ■ Dynamic View

- A “virtual table” created dynamically upon request by a user
- No data actually stored; instead data from base table made available to user
- Based on SQL SELECT statement on base tables or other views

## ■ Materialized View

- Copy or replication of data
- Data actually stored
- Must be refreshed periodically to match the corresponding base tables

# Dynamic vs Materialized Views

- Advantages of Dynamic Views
  - Contain most current base table data
  - Use little storage space
- Disadvantages of Dynamic Views
  - Use processing time each time view is referenced
  - May or may not be directly updateable

# Processing Multiple Tables–Joins

- **Join**—a relational operation that causes two or more tables with a common domain to be combined into a single table or view
- **Natural join (inner join)**—a join in which the joining condition is based on equality between values in the common columns
- **Outer join**—a join in which rows that do not have matching values in common columns are nonetheless included in the result table (as opposed to *inner* join, in which rows must have matching values in order to appear in the result table)
- **Union join**—includes all columns from each table in the join, and an instance for each row of each table

The common columns in joined tables are usually the primary key of the dominant table and the foreign key of the dependent table in 1:M relationships

# Pine Valley Furniture Company Customer and Order tables with pointers from customers to their orders

The image shows two database tables side-by-side. The left table, 'Order\_T', has columns: OrderID, OrderDate, and CustomerID. The right table, 'Customer\_T', has columns: CustomerID, CustomerName, CustomerAddress, CustomerCity, CustomerState, and CustomerPostalCode. Arrows indicate a one-to-many relationship from Customer\_T to Order\_T. Each customer in the Customer\_T table is linked to one or more orders in the Order\_T table by their CustomerID.

| OrderID | OrderDate  | CustomerID |
|---------|------------|------------|
| 1001    | 10/21/2010 | 1          |
| 1002    | 10/21/2010 | 8          |
| 1003    | 10/22/2010 | 15         |
| 1004    | 10/22/2010 | 5          |
| 1005    | 10/24/2010 | 3          |
| 1006    | 10/24/2010 | 2          |
| 1007    | 10/27/2010 | 11         |
| 1008    | 10/30/2010 | 12         |
| 1009    | 11/5/2010  | 4          |
| 1010    | 11/5/2010  | 1          |
| 0       |            | 0          |

| CustomerID | CustomerName             | CustomerAddress      | CustomerCity  | CustomerState | CustomerPostalCode |
|------------|--------------------------|----------------------|---------------|---------------|--------------------|
| 1          | Contemporary Casuals     | 1355 S Hines Blvd    | Gainesville   | FL            | 32601-2871         |
| 2          | Value Furniture          | 15145 S.W. 17th St.  | Plano         | TX            | 75094-7743         |
| 3          | Home Furnishings         | 1900 Allard Ave.     | Albany        | NY            | 12209-1125         |
| 4          | Eastern Furniture        | 1925 Beltline Rd.    | Carteret      | NJ            | 07008-3188         |
| 5          | Impressions              | 5585 Westcott Ct.    | Sacramento    | CA            | 94206-4056         |
| 6          | Furniture Gallery        | 325 Flatiron Dr.     | Boulder       | CO            | 80514-4432         |
| 7          | Period Furniture         | 394 Rainbow Dr.      | Seattle       | WA            | 97954-5589         |
| 8          | California Classics      | 816 Peach Rd.        | Santa Clara   | CA            | 96915-7754         |
| 9          | M and H Casual Furniture | 3709 First Street    | Clearwater    | FL            | 34620-2314         |
| 10         | Seminole Interiors       | 2400 Rocky Point Dr. | Seminole      | FL            | 34646-4423         |
| 11         | American Euro Lifestyles | 2424 Missouri Ave N. | Prospect Park | NJ            | 07508-5621         |
| 12         | Battle Creek Furniture   | 345 Capitol Ave. SW  | Battle Creek  | MI            | 49015-3401         |
| 13         | Heritage Furnishings     | 66789 College Ave.   | Carlisle      | PA            | 17013-8834         |
| 14         | Kaneohe Homes            | 112 Kiowai St.       | Kaneohe       | HI            | 96744-2537         |
| 15         | Mountain Scenes          | 4132 Main Street     | Ogden         | UT            | 84403-4432         |
| (New)      |                          |                      |               |               |                    |

These tables are used in queries that follow

# Natural Join Example

- For each customer who placed an order, what is the customer's name and order number?

Join involves multiple tables in FROM clause

```
SELECT Customer_T.CustomerID, CustomerName, OrderID  
FROM Customer_T INNER JOIN Order_T ON
```

```
Customer_T.CustomerID = Order_T.CustomerID;
```

ON clause performs the equality check for common columns of the two tables

Note: from Fig. 1, you see that only 10 Customers have links with orders.

➔ Only 10 rows will be returned from this INNER join.

# Outer Join Example

- List the customer name, ID number, and order number for all customers. Include customer information even for customers that do not have an order

```
SELECT Customer_T.CustomerID, CustomerName, OrderID  
FROM Customer_T LEFT OUTER JOIN Order_T  
ON Customer_T.CustomerID = Order_T.CustomerID;
```

LEFT OUTER JOIN syntax with ON causes customer data to appear even if there is no corresponding order data

Unlike INNER join, this will include customer rows with no matching order rows

## Results

Unlike  
INNER join,  
this will  
include  
customer  
rows with  
no  
matching  
order rows

| CUSTOMERID        | CUSTOMERNAME             | ORDERID |
|-------------------|--------------------------|---------|
| 1                 | Contemporary Casuals     | 1001    |
| 1                 | Contemporary Casuals     | 1010    |
| 2                 | Value Furniture          | 1006    |
| 3                 | Home Furnishings         | 1005    |
| 4                 | Eastern Furniture        | 1009    |
| 5                 | Impressions              | 1004    |
| 6                 | Furniture Gallery        |         |
| 7                 | Period Furniture         |         |
| 8                 | California Classics      | 1002    |
| 9                 | M & H Casual Furniture   |         |
| 10                | Seminole Interiors       |         |
| 11                | American Euro Lifestyles | 1007    |
| 12                | Battle Creek Furniture   | 1008    |
| 13                | Heritage Furnishings     |         |
| 14                | Kaneohe Homes            |         |
| 15                | Mountain Scenes          | 1003    |
| 16 rows selected. |                          |         |

# Multiple Table Join Example

- Assemble all information necessary to create an invoice for order number 1006

Four tables involved in this join

```
SELECT Customer_T.CustomerID, CustomerName, CustomerAddress,  
       CustomerCity, CustomerState, CustomerPostalCode, Order_T.OrderID,  
       OrderDate, OrderedQuantity, ProductDescription, ProductStandardPrice,  
       (OrderedQuantity * ProductStandardPrice)
```

```
FROM Customer_T, Order_T, OrderLine_T, Product_T
```

```
WHERE Customer_T.CustomerID = Order_T.CustomerID  
      AND Order_T.OrderID = OrderLine_T.OrderID  
      AND OrderLine_T.ProductID = Product_T.ProductID  
      AND Order_T.OrderID = 1006;
```

Each pair of tables requires an equality-check condition in the WHERE clause, matching primary keys against foreign keys



# Self-Join Example

*Query:* What are the employee ID and name of each employee and the name of his or her supervisor (label the supervisor's name Manager)?

```
SELECT E.EmployeeID, E.EmployeeName, M.EmployeeName AS Manager  
FROM Employee_T E, Employee_T M  
WHERE E.EmployeeSupervisor = M.EmployeeID;
```

The same table is used on both sides of the join; distinguished using table aliases

*Result:*

| EMPLOYEEID | EMPLOYEENAME | MANAGER      |
|------------|--------------|--------------|
| 123-44-347 | Jim Jason    | Robert Lewis |

Self-joins are usually used on tables with unary relationships

# Processing Multiple Tables Using Subqueries

- Subquery—placing an inner query (SELECT statement) inside an outer query
- Options:
  - In a condition of the WHERE clause
  - As a “table” of the FROM clause
  - Within the HAVING clause
- Subqueries can be:
  - Noncorrelated—executed once for the entire outer query
  - Correlated—executed once for each row returned by the outer query

# Subquery Example

- Show all customers who have placed an order

The IN operator will test to see if the CUSTOMER\_ID value of a row is included in the list returned from the subquery

```
SELECT CustomerName FROM Customer_T  
WHERE CustomerID IN
```

```
(SELECT DISTINCT CustomerID FROM Order_T);
```

Subquery is embedded in parentheses. In this case it returns a list that will be used in the WHERE clause of the outer query

Result:

| <u>CUSTOMER_NAME</u>     |
|--------------------------|
| Contemporary Casuals     |
| Value Furniture          |
| Home Furnishings         |
| Eastern Furniture        |
| Impressions              |
| California Classics      |
| American Euro Lifestyles |
| Battle Creek Furniture   |
| Mountain Scenes          |
| 9 rows selected.         |

# Processing a noncorrelated subquery

1. The subquery executes and returns the customer IDs from the ORDER\_T table
2. The outer query on the results of the subquery

```
SELECT CUSTOMER_NAME  
FROM CUSTOMER_T  
WHERE CUSTOMER_ID IN
```

(SELECT DISTINCT CUSTOMER\_ID  
FROM ORDER\_T);

1. The subquery (shown in the box) is processed first and an intermediate results table created:

| CUSTOMER_ID |
|-------------|
| 1           |
| 8           |
| 15          |
| 5           |
| 3           |
| 2           |
| 11          |
| 12          |
| 4           |

9 rows selected.

No reference to data in outer query, so subquery executes once only

2. The outer query returns the requested customer information for each customer included in the intermediate results table:

| CUSTOMER_NAME            |
|--------------------------|
| Contemporary Casuals     |
| Value Furniture          |
| Home Furnishings         |
| Eastern Furniture        |
| Impressions              |
| California Classics      |
| American Euro Lifestyles |
| Battle Creek Furniture   |
| Mountain Scenes          |

9 rows selected.

These are the only customers that have IDs in the ORDER\_T table

# Correlated vs. Noncorrelated Subqueries

- Noncorrelated subqueries:
  - Do not depend on data from the outer query
  - Execute once for the entire outer query
- Correlated subqueries:
  - Make use of data from the outer query
  - Execute once for each row of the outer query
  - Can use the EXISTS operator

# Correlated Subquery Example

- Show all orders that include furniture finished in natural ash

The EXISTS operator will return a TRUE value if the subquery resulted in a non-empty set, otherwise it returns a FALSE

```
SELECT DISTINCT OrderID FROM OrderLine_T
WHERE EXISTS
  (SELECT * FROM Product_T
   WHERE ProductID = OrderLine_T.ProductID
   AND ProductFinish = 'Natural ash');
```

The subquery is testing for a value that comes from the outer query

# Processing a correlated subquery

```
SELECT DISTINCT ORDER_ID FROM ORDER_LINE_T
WHERE EXISTS
  (SELECT *
   FROM PRODUCT_T
    WHERE PRODUCT_ID = ORDER_LINE_T.PRODUCT_ID
      AND PRODUCT_FINISH = 'Natural Ash');
```

Subquery refers to outer-query data, so executes once for each row of outer query

| Order_ID | Product_ID | Ordered_Quantity |
|----------|------------|------------------|
| 1001     | 1          | 1                |
| 1001     | 4          | 1                |
| 1001     | 3          | 1                |
| 1001     | 3          | 1                |
| 1004     | 6          | 1                |
| 1004     | 6          | 1                |
| 1005     | 4          | 1                |
| 1005     | 4          | 1                |
| 1005     | 6          | 1                |
| 1007     | 1          | 1                |
| 1007     | 2          | 1                |
| 1008     | 3          | 1                |
| 1008     | 6          | 1                |
| 1008     | 4          | 1                |
| 1008     | 7          | 1                |
| 1010     | 8          | 10               |
| 1010     | 8          | 10               |

|   | Product_ID   | Product Description  | Product_Finish | Standard_Price | Product_Line_Id |
|---|--------------|----------------------|----------------|----------------|-----------------|
| 1 | 1            | End Table            | Cherry         | \$175.00       | 10001           |
| 2 | 2            | Coffee Table         | Natural Ash    | \$200.00       | 20001           |
| 3 | 3            | Computer Desk        | Natural Ash    | \$375.00       | 20001           |
| 4 | 4            | Entertainment Center | Natural Maple  | \$650.00       | 30001           |
| 5 | 5            | Writer's Desk        | Cherry         | \$325.00       | 10001           |
| 6 | 6            | 8-Drawer Dresser     | White Ash      | \$750.00       | 20001           |
| 7 | 7            | Dining Table         | Natural Ash    | \$800.00       | 20001           |
| 8 | 8            | Computer Desk        | Walnut         | \$250.00       | 30001           |
| * | (AutoNumber) |                      |                | \$0.00         |                 |

1. The first order ID is selected from ORDER\_LINE\_T: ORDER\_ID =1001.
2. The subquery is evaluated to see if any product in that order has a natural ash finish. Product 2 does, and is part of the order. EXISTS is valued as true and the order ID is added to the result table.
3. The next order ID is selected from ORDER\_LINE\_T: ORDER\_ID =1002.
4. The subquery is evaluated to see if the product ordered has a natural ash finish. It does. EXISTS is valued as true and the order ID is added to the result table.
5. Processing continues through each order ID. Orders 1004, 1005, and 1010 are not included in the result table because they do not include any furniture with a natural ash finish. The final result table is shown in the text on page 303.

Note: only the orders that involve products with Natural Ash will be included in the final results

# Another Subquery Example

- Show all products whose standard price is higher than the average price

Subquery forms the derived table used in the FROM clause of the outer query

One column of the subquery is an aggregate function that has an alias name. That alias can then be referred to in the outer query

```
SELECT ProductDescription, ProductStandardPrice, AvgPrice  
FROM
```

```
(SELECT AVG(ProductStandardPrice) AS AvgPrice FROM Product_T) AP,  
Product_T  
WHERE ProductStandardPrice > AvgPrice;
```

The WHERE clause normally cannot include aggregate functions, but because the aggregate is performed in the subquery its result can be used in the outer query's WHERE clause

AP is the name of the result table of the subquery.



# Union Queries

- Combine the output (union of multiple queries) together into a single result table

```
SELECT C1.CUSTOMER_ID,CUSTOMER_NAME,ORDERED_QUANTITY,  
'Largest Quantity' QUANTITY  
FROM CUSTOMER_T C1,ORDER_T O1, ORDER_LINE_T Q1  
WHERE C1.CUSTOMER_ID =O1.CUSTOMER_ID  
AND O1.ORDER_ID =Q1.ORDER_ID  
AND ORDERED_QUANTITY =  
      (SELECT MAX(ORDERED_QUANTITY)  
       FROM ORDER_LINE_T)
```

First query

Combine → **UNION**

```
SELECT C1.CUSTOMER_ID,CUSTOMER_NAME,ORDERED_QUANTITY,  
'Smallest Quantity'  
FROM CUSTOMER_T C1,ORDER_T O1, ORDER_LINE_T Q1  
WHERE C1.CUSTOMER_ID =O1.CUSTOMER_ID  
AND O1.ORDER_ID =Q1.ORDER_ID  
AND ORDERED_QUANTITY =  
      (SELECT MIN(ORDERED_QUANTITY)  
       FROM ORDER_LINE_T)  
ORDER BY ORDERED_QUANTITY;
```

Second query

# Conditional Expressions Using Case Syntax

This is available with newer versions of SQL, previously not part of the standard

```
{CASE expression  
{WHEN expression  
THEN {expression | NULL}} ...  
| {WHEN predicate  
THEN {expression | NULL}} ...  
[ELSE {expression | NULL}]  
END }  
| ( NULLIF (expression, expression) )  
| ( COALESCE (expression . . . ) }
```

```
SELECT CASE  
    WHEN ProductLine = 1 THEN ProductDescription  
    ELSE '####'  
END AS ProductDescription  
FROM Product_T;
```

# Tips for Developing Queries

- Be familiar with the data model (entities and relationships)
- Understand the desired results
- Know the attributes desired in result
- Identify the entities that contain desired attributes
- Review ERD
- Construct a WHERE equality for each link
- Fine tune with GROUP BY and HAVING clauses if needed
- Consider the effect on unusual data

# Query Efficiency Considerations

- Instead of `SELECT *`, identify the specific attributes in the `SELECT` clause; this helps reduce network traffic of result set
- Limit the number of subqueries; try to make everything done in a single query if possible
- If data is to be used many times, make a separate query and store its results rather than performing the query repeatedly

# Guidelines for Better Query Design

- Understand how indexes are used in query processing
- Keep optimizer statistics up-to-date
- Use compatible data types for fields and literals
- Write simple queries
- Break complex queries into multiple simple parts
- Don't nest one query inside another query
- Don't combine a query with itself (if possible avoid self-joins)
- Create temporary tables for groups of queries
- Combine update operations
- Retrieve only the data you need
- Don't have the DBMS sort without an index
- Consider the total query processing time for ad hoc queries

# Ensuring Transaction Integrity

- Transaction = A discrete unit of work that must be completely processed or not processed at all
  - May involve multiple updates
  - If any update fails, then all other updates must be cancelled
- SQL commands for transactions
  - BEGIN TRANSACTION/END TRANSACTION
    - Marks boundaries of a transaction
  - COMMIT
    - Makes all updates permanent
  - ROLLBACK
    - Cancels updates since the last COMMIT

# An SQL Transaction sequence (in pseudocode)

BEGIN transaction

INSERT OrderID, Orderdate, CustomerID into Order\_T;

INSERT OrderID, ProductID, OrderedQuantity into OrderLine\_T;

INSERT OrderID, ProductID, OrderedQuantity into OrderLine\_T;

INSERT OrderID, ProductID, OrderedQuantity into OrderLine\_T;

END transaction

Valid information inserted.  
COMMIT work.



All changes to data  
are made permanent.

Invalid ProductID entered.



Transaction will be ABORTED.  
ROLLBACK all changes made to Order\_T.



All changes made to Order\_T  
and OrderLine\_T are removed.  
Database state is just as it was  
before the transaction began.

# Data Dictionary Facilities

- System tables that store metadata
- Users usually can view some of these tables
- Users are restricted from updating them
- Some examples in Oracle 10g/11g
  - DBA\_TABLES – descriptions of tables
  - DBA\_CONSTRAINTS – description of constraints
  - DBA\_USERS – information about the users of the system
- Examples in Microsoft SQL Server 2008
  - sys.columns – table and column definitions
  - sys.indexes – table index information
  - sys.foreign\_key\_columns – details about columns in foreign key constraints



# Enhancements/Extensions in Newer Standards

- User-defined data types (UDT)
  - Subclasses of standard types or an object type
- Analytical functions (for OLAP)
  - CEILING, FLOOR, SQRT, RANK, DENSE\_RANK
  - WINDOW—improved numerical analysis capabilities
- New Data Types
  - BIGINT, MULTISSET (collection), XML
- CREATE TABLE LIKE—create a new table similar to an existing one
- MERGE

# Merge Statement

```
MERGE INTO Product_T AS PROD
USING
(SELECT ProductID, ProductDescription, ProductFinish,
ProductStandardPrice, ProductLineID FROM Purchases_T) AS PURCH
  ON (PROD.ProductID = PURCH.ProductID)
WHEN MATCHED THEN UPDATE
  PROD.ProductStandardPrice = PURCH.ProductStandardPrice
WHEN NOT MATCHED THEN INSERT
  (ProductID, ProductDescription, ProductFinish, ProductStandardPrice,
  ProductLineID)
VALUES(PURCH.ProductID, PURCH.ProductDescription,
PURCH.ProductFinish, PURCH.ProductStandardPrice,
PURCH.ProductLineID);
```

Makes it easier to update a table...allows combination of Insert and Update in one statement

Useful for updating master tables with new data

# Enhancements/Extensions in Newer Standards

- Persistent Stored Modules (SQL/PSM)
  - Capability to create and drop code modules
  - New statements:
    - CASE, IF, LOOP, FOR, WHILE, etc.
    - Makes SQL into a procedural language
- Oracle has proprietary version called PL/SQL, and Microsoft SQL Server has Transact/SQL

# Routines and Triggers

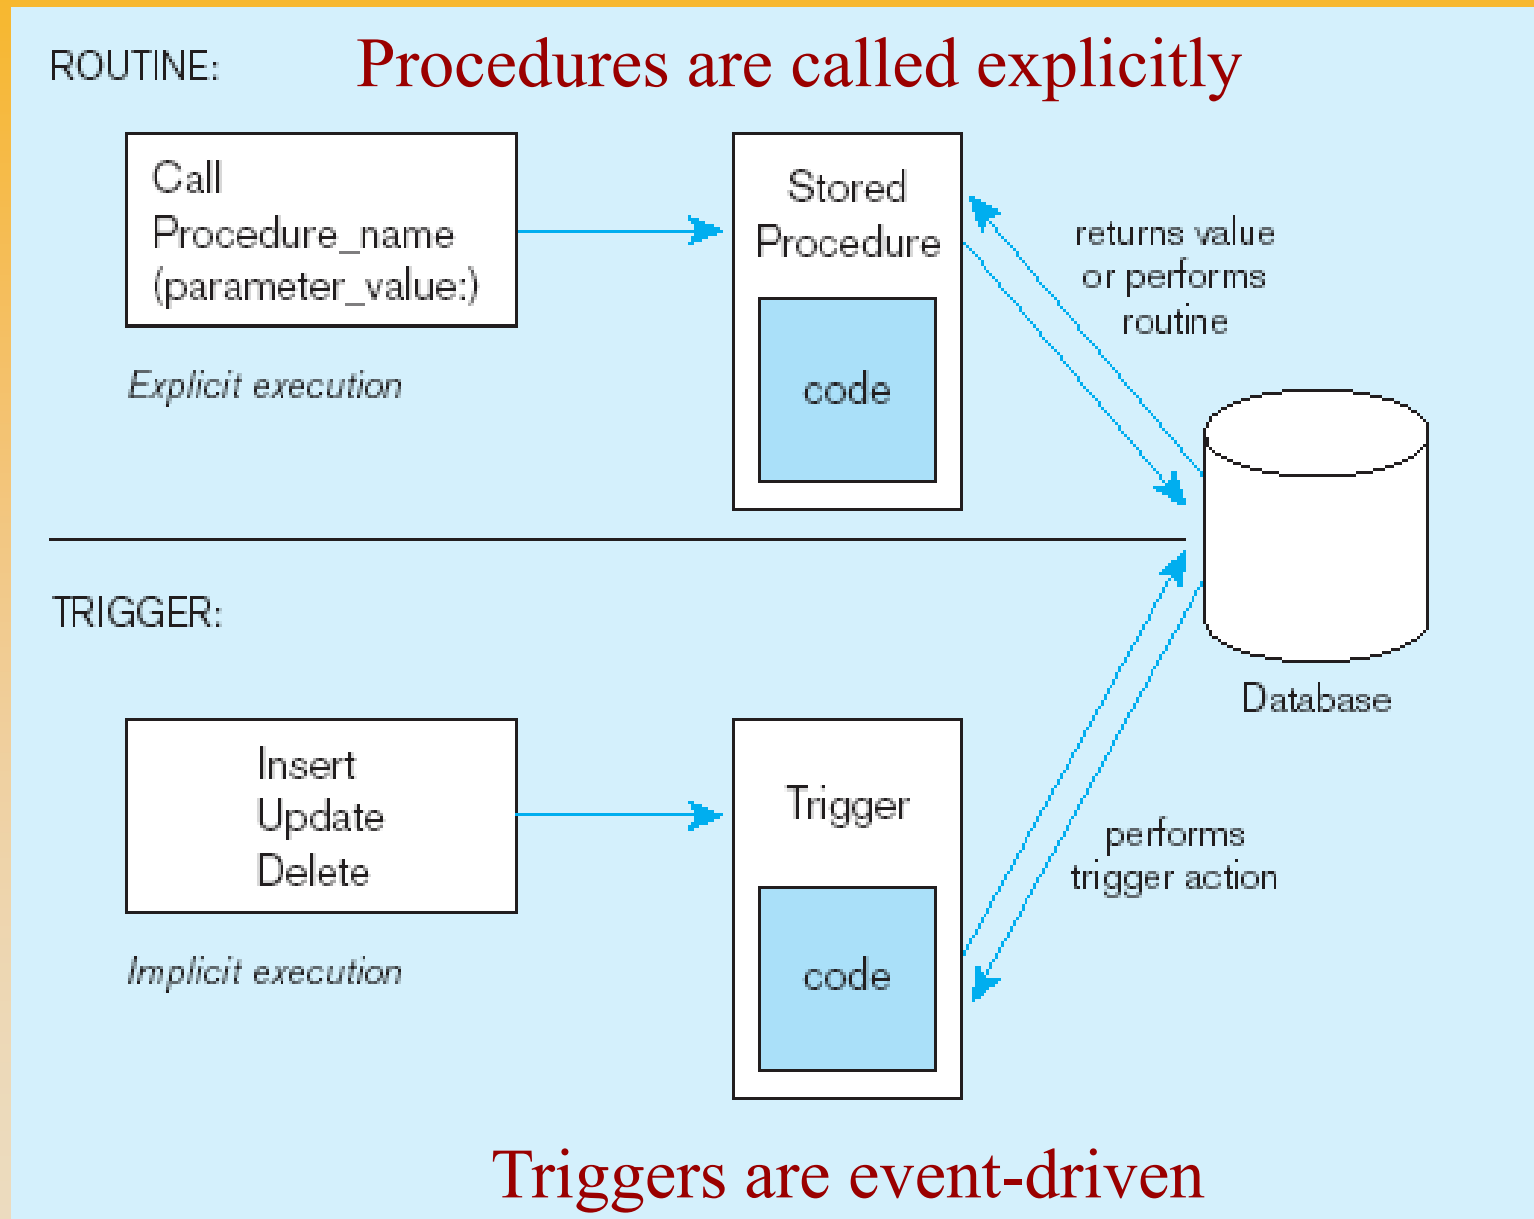
## ■ **Routines**

- Program modules that execute on demand
- **Functions**—routines that return values and take input parameters
- **Procedures**—routines that do not return values and can take input or output parameters

## ■ **Triggers**

- Routines that execute in response to a database event (INSERT, UPDATE, or DELETE)

# Triggers contrasted with stored procedures



Source: adapted from Mullins, 1995.

## Simplified trigger syntax

```
CREATE TRIGGER trigger_name
    {BEFORE | AFTER | INSTEAD OF} {INSERT | DELETE | UPDATE} ON
    table_name
    [FOR EACH {ROW | STATEMENT}] [WHEN (search condition)]
    <triggered SQL statement here>;
```

## Create routine syntax

```
{CREATE PROCEDURE | CREATE FUNCTION} routine_name
([parameter [{parameter} . . .]])
[RETURNS data_type result_cast] /* for functions only */
[LANGUAGE {ADA | C | COBOL | FORTRAN | MUMPS | PASCAL | PLI | SQL}]
[PARAMETER STYLE {SQL | GENERAL}]
[SPECIFIC specific_name]
[DETERMINISTIC | NOT DETERMINISTIC]
[NO SQL | CONTAINS SQL | READS SQL DATA | MODIFIES SQL DATA]
[RETURNS NULL ON NULL INPUT | CALLED ON NULL INPUT]
[DYNAMIC RESULT SETS unsigned_integer] /* for procedures only */
[STATIC DISPATCH] /* for functions only */
[NEW SAVEPOINT LEVEL | OLD SAVEPOINT LEVEL]
routine_body
```