

Introduction to NoSQL Databases

MSBA7024 / MACC7020

Database Design and Management

Objectives

- Understand the shortcomings of RDBMS in today's applications
- Describe the basics of NoSQL and why it has become popular
- Describe the major properties of NoSQL databases
- Compare the pros and cons of RDBMS and NoSQL models
- Discuss and compare the different types of NoSQL models
- Understand the CAP theorem

Background

- Relational database management systems (RDBMS)
 - mainstay of business
 - sometimes also called SQL databases
- Web-based applications caused spikes
 - explosion of social media sites (Facebook, Twitter) with large data needs
 - rise of cloud-based solutions such as Amazon S3 (simple storage solution)
- Hooking RDBMS to web-based application becomes troublesome


Characteristics of RDBMS

- RDBMSs have been around for ages
 - MySQL is one of the most popular among them
- Data stored in tables
- Schema-based, i.e., structured tables
- Each row (data item) in a table has a primary key that is unique within that table
- Queried using SQL (Structured Query Language)
- Support joins

Relational Database Example

users table

user_id	name	zipcode	blog_url	blog_id
101	Alice	12345	alice.net	1
422	Charlie	45783	charlie.com	3
555	Bob	99910	bob.blogspot.com	2



blog table

id	url	last_updated	num_posts
1	alice.net	5/2/14	332
2	bob.blogspot.com	4/2/13	10003
3	charlie.com	6/15/14	7

Example SQL queries

1. `SELECT zipcode
FROM users
WHERE name = "Bob"`
2. `SELECT url
FROM blog
WHERE id = 3`
3. `SELECT users.zipcode,
blog.num_posts
FROM users JOIN blog
ON users.blog_url = blog.url`

Popular RDBMS



Mismatch with Today's Workloads

- Data: Large and unstructured
- Lots of random reads and writes
- Sometimes write-heavy
- Foreign keys less needed
- Joins infrequent

Needs of Today's Workloads

- Speed
- Avoid Single point of Failure (SPoF)
- Low TCO (Total cost of operation)
- Fewer system administrators
- Incremental Scalability
- Scale out, not up

Scale out, not Scale up

- Scale up = grow your cluster capacity by replacing with more powerful machines
 - Traditional approach
 - a.k.a. vertical scaling: make a “single” machine more powerful
 - Not cost-effective
 - And you need to replace machines often
- Scale out = incrementally grow your cluster capacity by adding more COTS machines (Components Off the Shelf)
 - Cheaper
 - a.k.a. horizontal scaling
 - Over a long duration, phase in a few newer (faster) machines as you phase out a few older machines
 - Used by most companies who run datacenters and clouds today

Scaling out RDBMS: Master/Slave

- Master/Slave
 - All writes are written to the master
 - All reads performed against the replicated slave databases
 - Critical reads may be incorrect as writes may not have been propagated down
 - Large datasets can pose problems as master needs to duplicate data to slaves

Scaling out RDBMS: Sharding

- Sharding (Partitioning)
 - Scales well for both reads and writes
 - Not transparent, application needs to be partition-aware
 - Can no longer have relationships/joins across partitions
 - Loss of referential integrity across shards

RDBMS

- ❑ Relational databases were not built for **distributed applications**.


Because...

- ❑ Joins are expensive
- ❑ Hard to scale horizontally
- ❑ Impedance mismatch occurs
- ❑ Expensive (product cost, hardware, Maintenance)

And....

It's weak in:

- ❑ Speed (performance)
- ❑ Availability
- ❑ Partition tolerance



NoSQL Data Model

- NoSQL = “Not Only SQL”
- Necessary API operations: `get(key)` and `put(key, value)`
 - And some extended operations, e.g., “CQL” in Cassandra key-value store
- Tables
 - “Column families” in Cassandra, “Table” in HBase, “Collection” in MongoDB
 - Like RDBMS tables, but:
 - May be unstructured: May not have schemas
 - Some columns may be missing from some rows
 - Do not always support joins or have foreign keys
 - Can have index tables, just like RDBMS

NoSQL Data Model

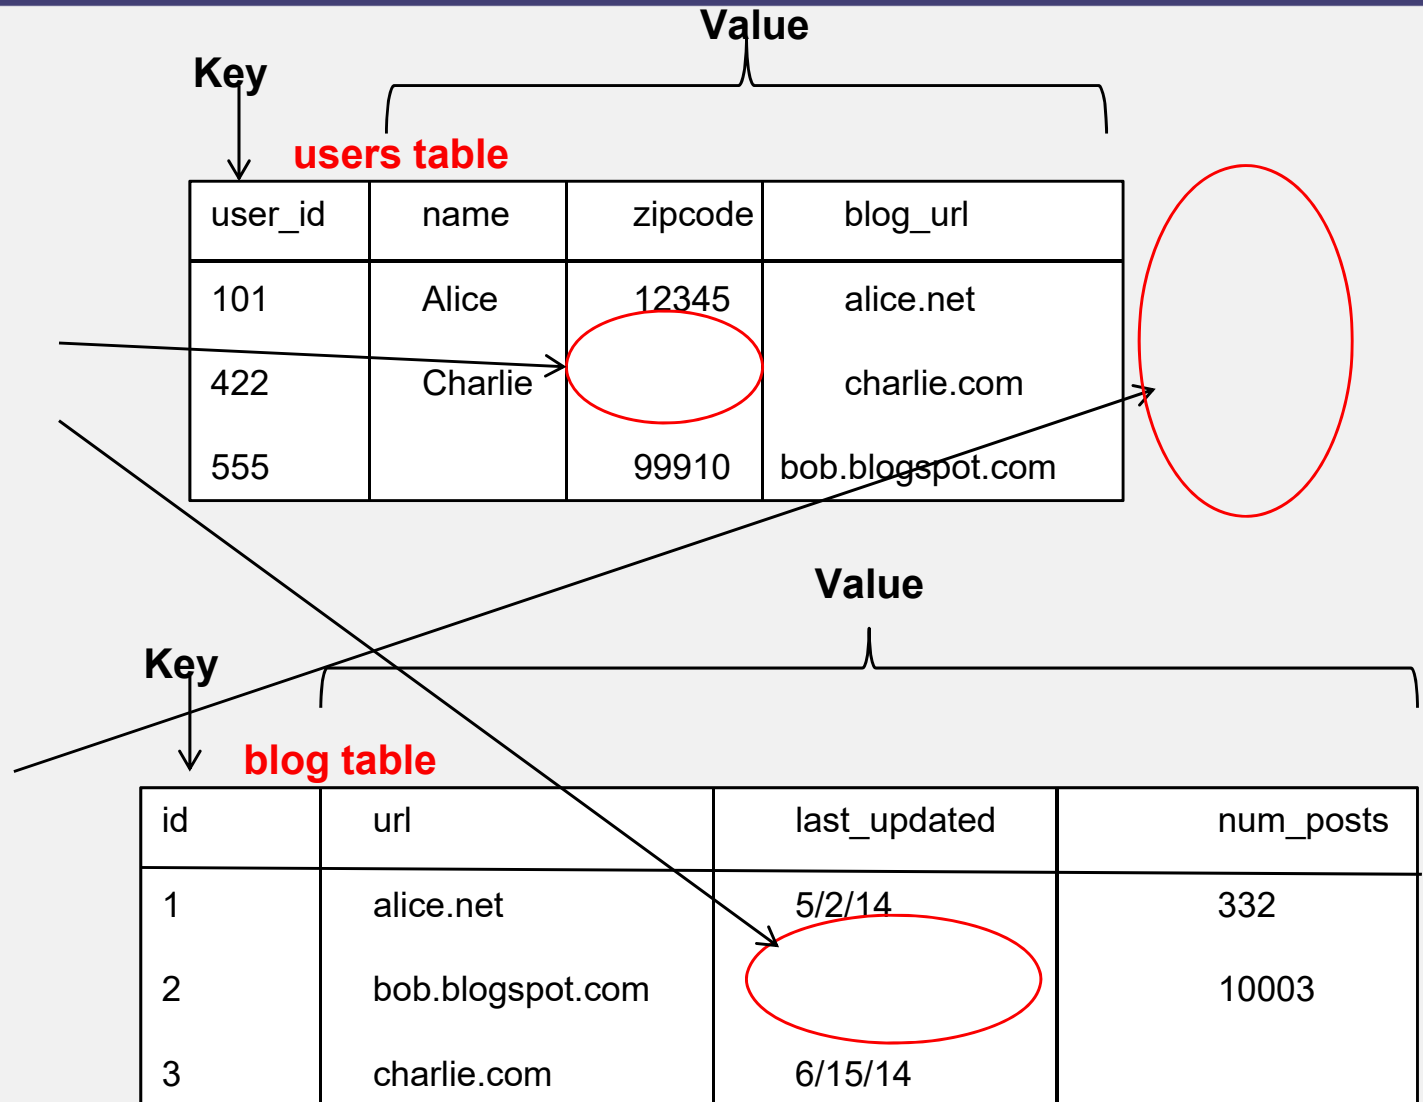
- Key features (advantages):
 - Non-relational
 - Do not require schema
 - Data are replicated to multiple nodes and can be partitioned:
 - down nodes easily replaced
 - no single point of failure
 - Horizontal scalable
 - Cheap, easy to implement (open-source)
 - Massive write performance
 - Fast key-value access

NoSQL Data Model

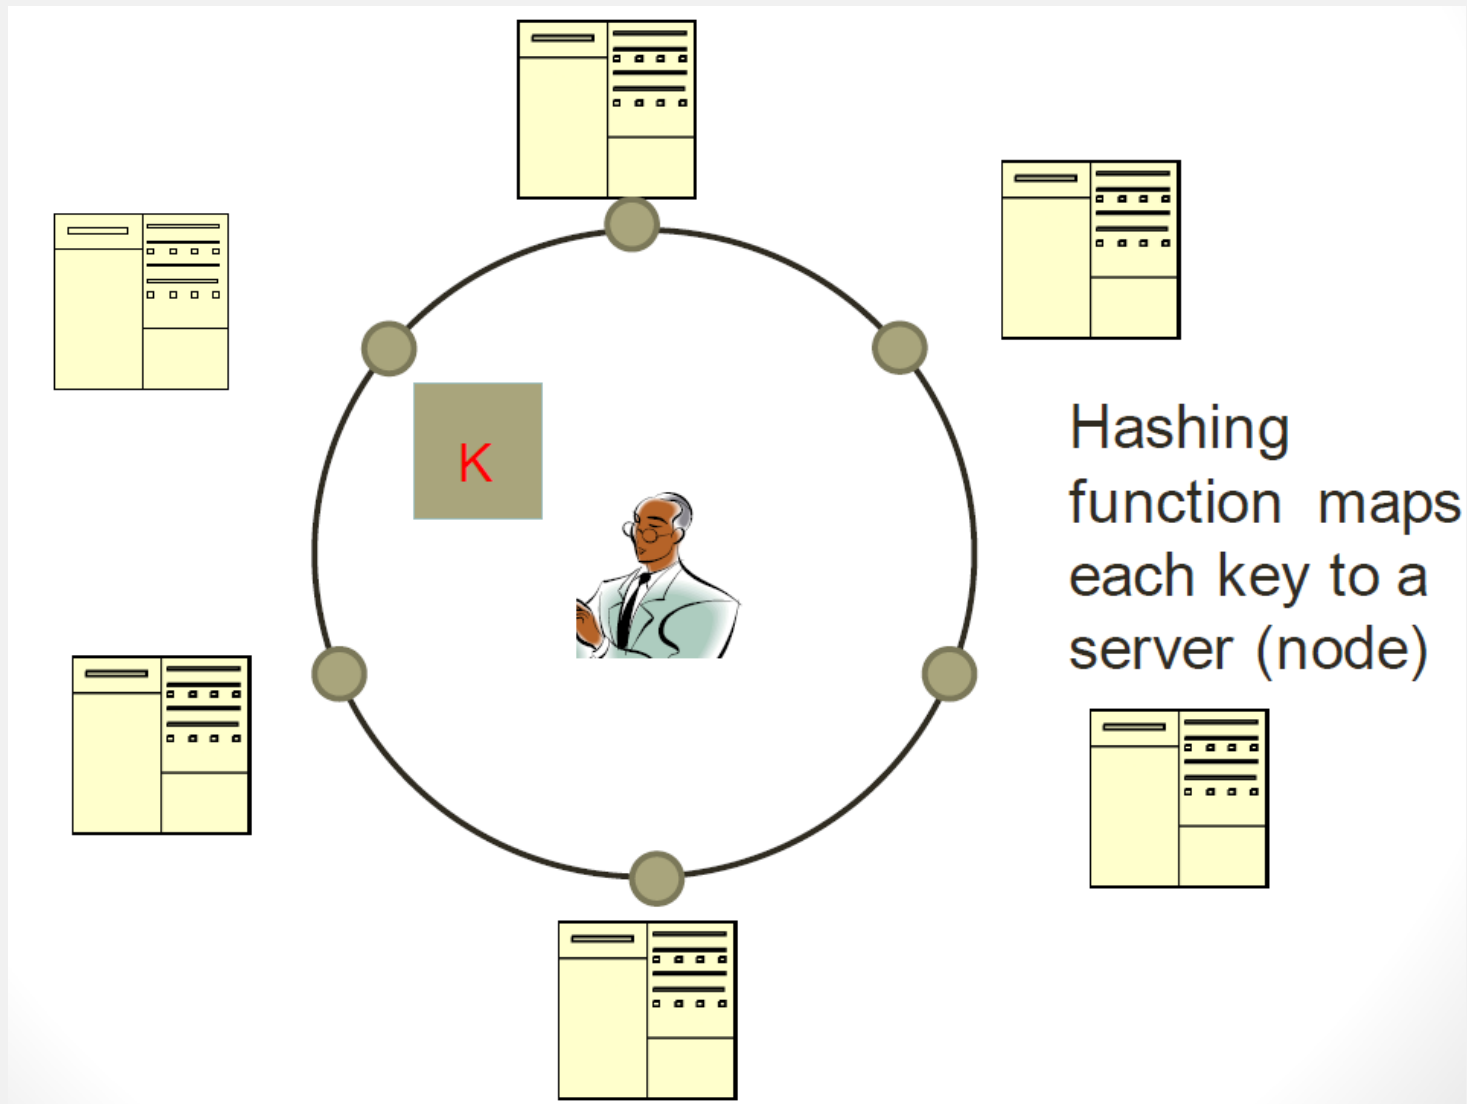


NoSQL Data Model

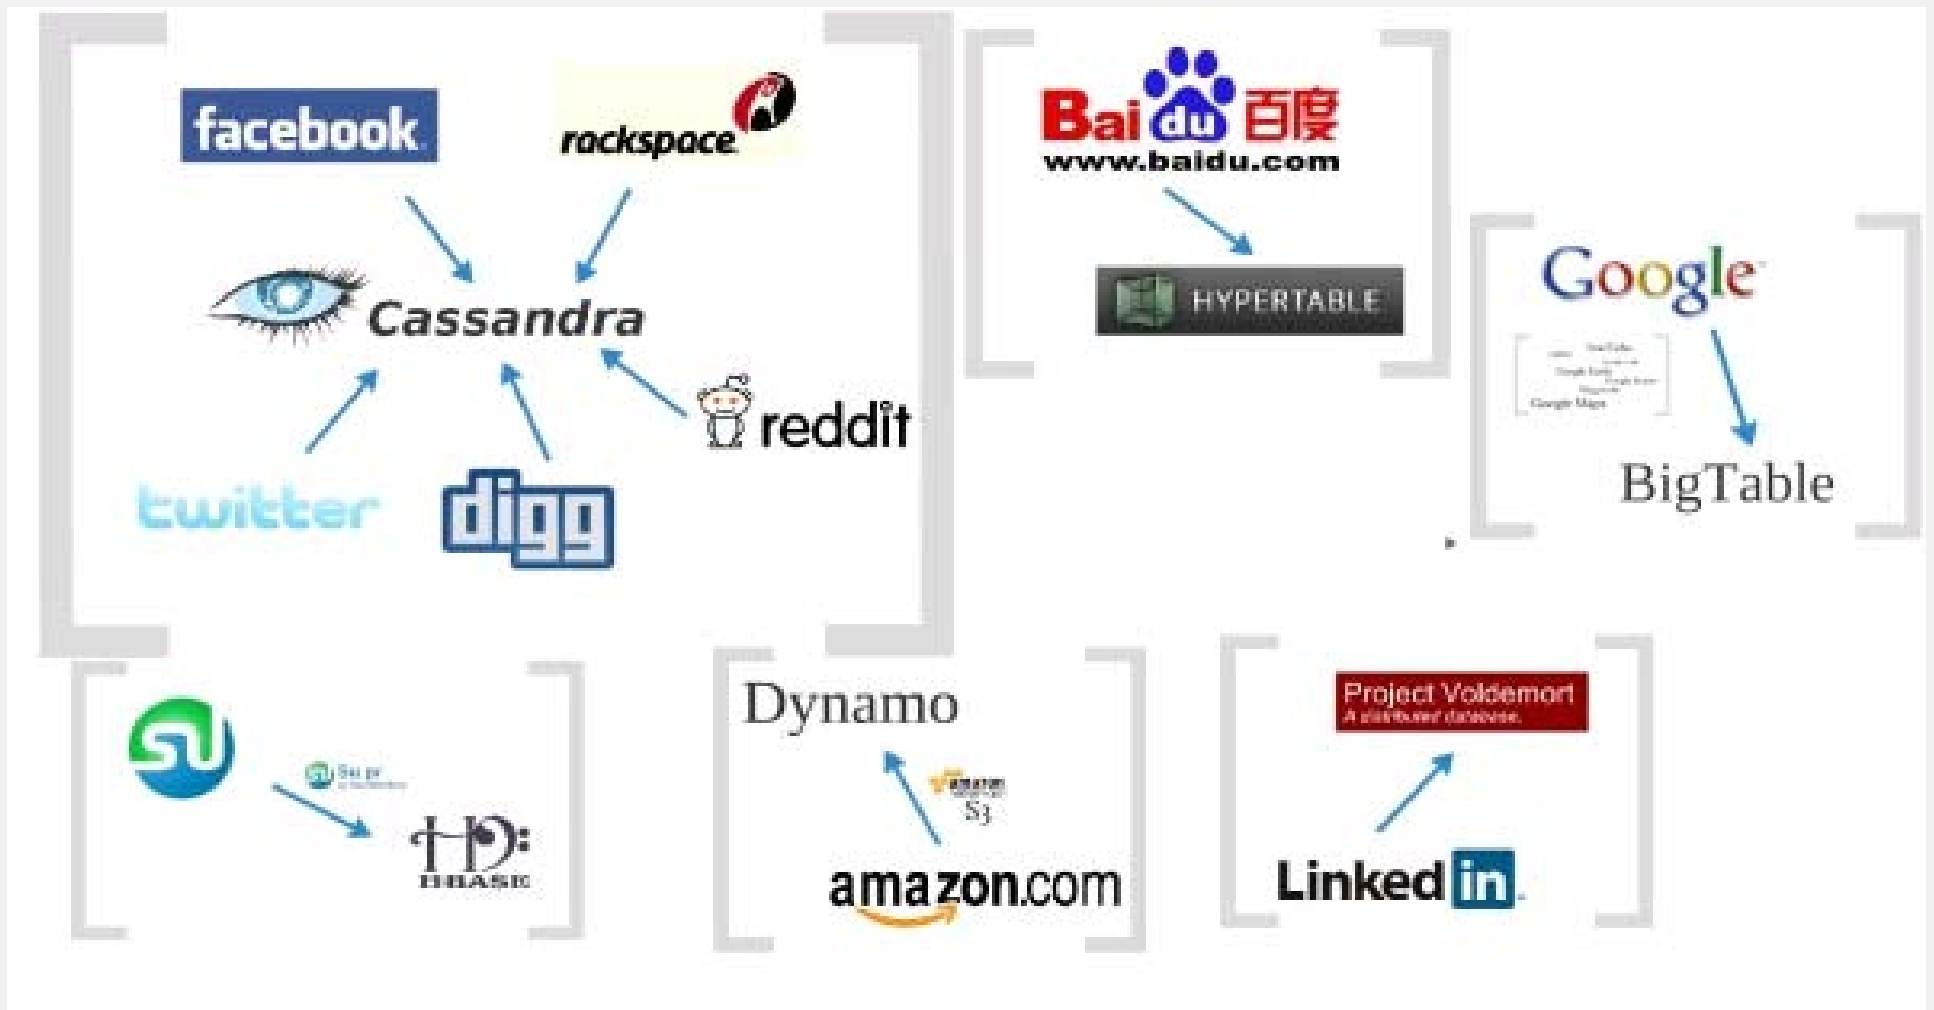
- Unstructured
- Columns Missing from some Rows
- No schema imposed
- No foreign keys, joins may not be supported



Typical NoSQL Architecture



Who are using them?



RDBMS vs. NoSQL

- RDBMS provide **ACID**



RDBMS vs. NoSQL

- NoSQL databases like MongoDB and Cassandra provide **BASE**
 - Basically Available Soft-state Eventual Consistency

Benefits of NoSQL

- **Elastic Scaling**
 - RDBMS scale up – bigger load , bigger server
 - NoSQL scale out – distribute data across multiple hosts seamlessly
- **DBA Specialists**
 - RDBMS require highly trained expert to monitor DB
 - NoSQL require less management, automatic repair and simpler data models
- **Big Data**
 - Huge increase in data RDBMS: capacity and constraints of data volumes at its limits
 - NoSQL designed for big data

Benefits of NoSQL

- **Flexible data models**
 - Change management to schema for RDBMS have to be carefully managed
 - NoSQL databases more relaxed in structure of data
 - Database schema changes do not have to be managed as one complicated change unit
 - Application already written to address an amorphous schema
- **Economics**
 - RDBMS rely on expensive proprietary servers to manage data
 - NoSQL: clusters of cheap commodity servers to manage the data and transaction volumes
 - Cost per gigabyte or transaction/second for NoSQL can be lower than the cost for a RDBMS

Drawbacks of NoSQL

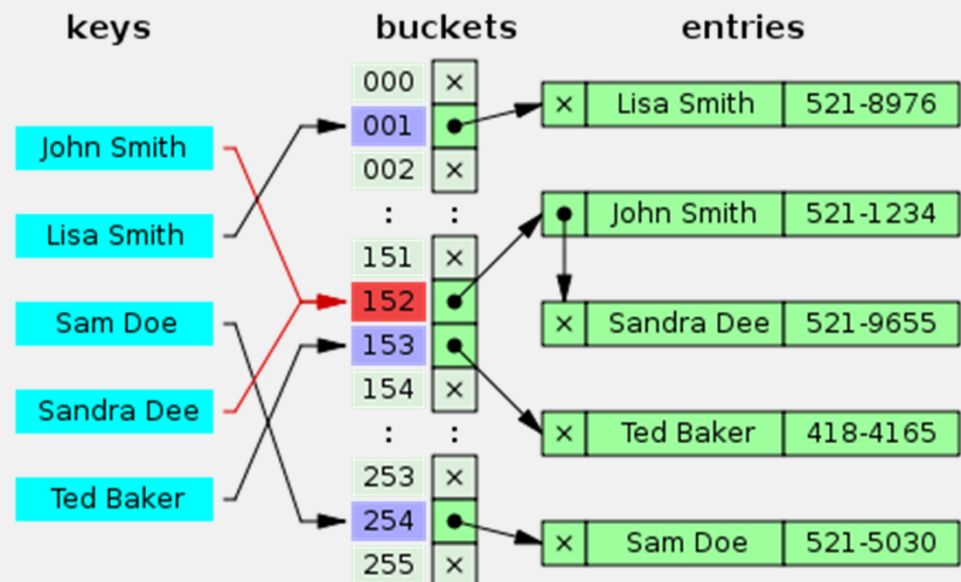
- **Support**
 - RDBMS vendors provide a high level of support to clients
 - Stellar reputation
 - NoSQL are open source projects with startups supporting them
 - Reputation not yet established
- **Maturity**
 - RDBMS mature product: means stable and dependable
 - Also means old no longer cutting edge nor interesting
 - NoSQL are still implementing their basic feature set

Drawbacks of NoSQL

- **Administration**
 - RDBMS administrator has a well-defined role
 - NoSQL's goal: no administrator necessary; however, NoSQL still requires effort to maintain
- **Lack of Expertise**
 - Whole workforce of trained and seasoned RDBMS developers
 - Still recruiting developers to the NoSQL camp
- **Analytics and Business Intelligence**
 - RDBMS designed to support decision-making
 - Data warehouse
 - NoSQL designed to meet the needs of Web 2.0 applications—not designed for ad hoc query of the data
 - Tools are being developed to address this need
 - More flexible to include new and unstructured data


Key-value Store

- NoSQL databases generally rely on key-value store.
- It is a dictionary data structure
 - Insert, lookup, and delete by key
 - Usually based on hash
 - But distributed



Key-value Store

- Also called a Distributed Hash Table (DHT)
- Main idea: partition set of key-values across many machines



Key-value Store

- Business: Key \rightarrow Value
- twitter.com: tweet id \rightarrow information about tweet
- amazon.com: item number \rightarrow information about it
- facebook.com: user id \rightarrow user profile, photos, etc.
- kayak.com: flight number \rightarrow information about flight, e.g., availability
- yourbank.com: account number \rightarrow account balances, transaction histories

Key-value Store

- Basic access:
 - `get(key)`: extract the value given a key
 - `put(key, value)`: create or update the value given its key
 - `delete(key)`: remove the key and its associated value
 - `execute(key, operation, parameters)`: invoke an operation to the value (given its key) which is a special data structure (e.g. List, Set, Map etc.)


Key-value Store

- Can handle large volumes of data, e.g., PBs
 - Can distribute data over hundreds, even thousands of machines
 - Designed to be faster with lower overhead (additional storage) than conventional DBMS.

NoSQL categories

1. Standard key-value
 - Example: DynamoDB, Voldermort, Scalaris
2. Document-based
 - Example: MongoDB, CouchDB
3. Column-based
 - Example: BigTable, Cassandra, HBase
4. Graph-based
 - Example: Neo4J, InfoGrid
5. Vector-based
 - Example: Pinecone, Milvus, Chroma
 - “No-schema” is a common characteristics of most NoSQL storage systems
 - Provide “flexible” data types

Key-value



- Focus on scaling to huge amounts of data
- Designed to handle massive load
- Based on Amazon's dynamo paper
- Data model: (global) collection of Key-value pairs
- *Dynamo ring partitioning and replication*
- Example: (DynamoDB)
 - *items* having one or more attributes (name, value)
 - An *attribute* can be single-valued or multi-valued like set.
 - items are combined into a *table*

Key-value

Pros:

- Very fast
- Very scalable (horizontally distributed to nodes based on key)
- Simple data model
- Eventual consistency
- Fault-tolerance

Cons:


- Simple key-value cannot model more complex data structure such as objects

Key-value

Name	Producer	Data model	Querying
SimpleDB	Amazon	set of couples (key, {attribute}), where attribute is a couple (name, value)	restricted SQL; select, delete, GetAttributes, and PutAttributes operations
Redis	Salvatore Sanfilippo	set of couples (key, value), where value is simple typed value, list, ordered (according to ranking) or unordered set, hash value	primitive operations for each value type
Dynamo	Amazon	like SimpleDB	simple get operation and put in a context
Voldemort	LinkedIn	like SimpleDB	similar to Dynamo

Document-based

- Can model more complex objects
- Inspired by Lotus Notes
- Data model: collection of documents
- Document: JSON (**J**ava**S**cript **O**bject **N**otation)
 - A data model, key-value pairs, which supports objects, records, structs, lists, array, maps, dates, Boolean, XML, and other semi-structured formats.

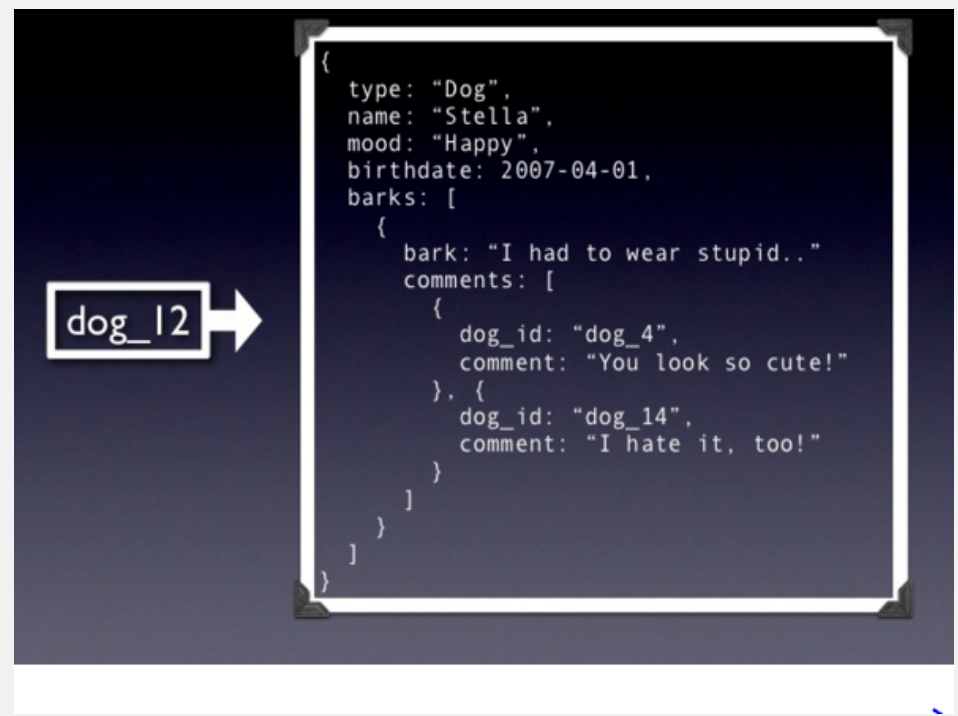
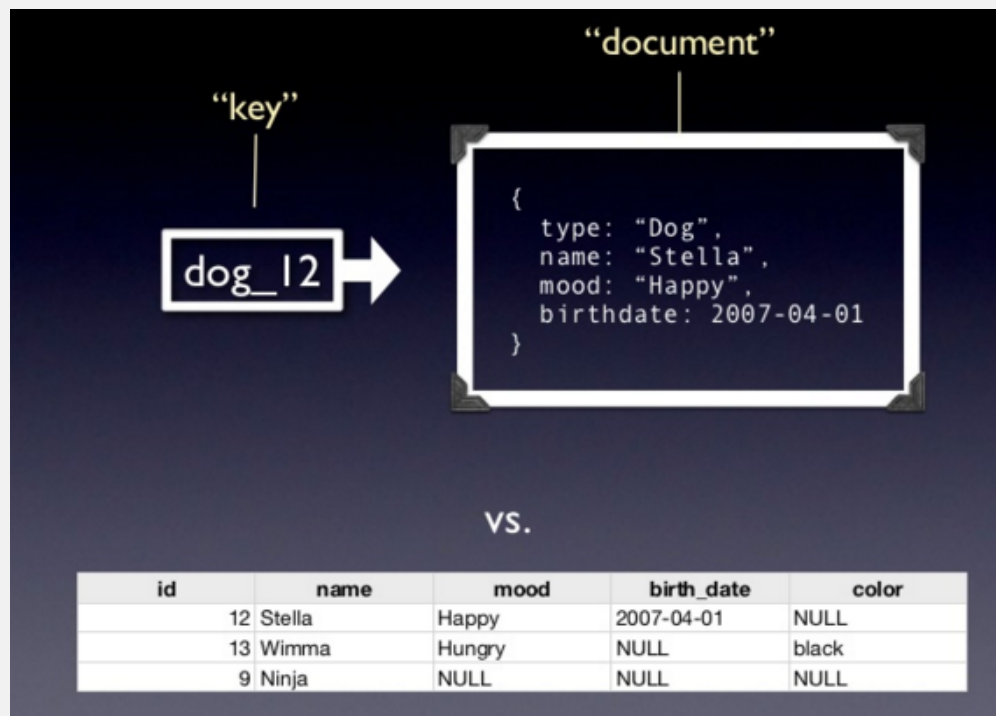


Document-based

- Example: (MongoDB) document

```
{  
  Name: "Jaroslav",  
  Address: "Malostranske nám. 25, 118 00 Praha 1",  
  Grandchildren: {Claire: "7", Barbara: "6", "Magda":  
    "3", "Kirsten: "1", "Otis: "3", Richard: "1"}},  
  Phones: [ "123-456-7890", "234-567-8963" ]  
}
```

Document-based



Document-based

Name	Producer	Data model	Querying
MongoDB	10gen	object-structured documents stored in collections; each object has a primary key called ObjectId	manipulations with objects in collections (find object or objects via simple selections and logical expressions, delete, update,)
Couchbase	Couchbase	document as a list of named (structured) items (JSON document)	by key and key range, views via Javascript and MapReduce


Column-based

- Based on Google's BigTable
- Google uses the key-value paradigm to map URLs to multidimensional data, such as:
 - Timestamps/Versions
 - Rank
 - Keywords
 - Links
- No explicit ordering is needed on keys since a hash function is used

Column-based

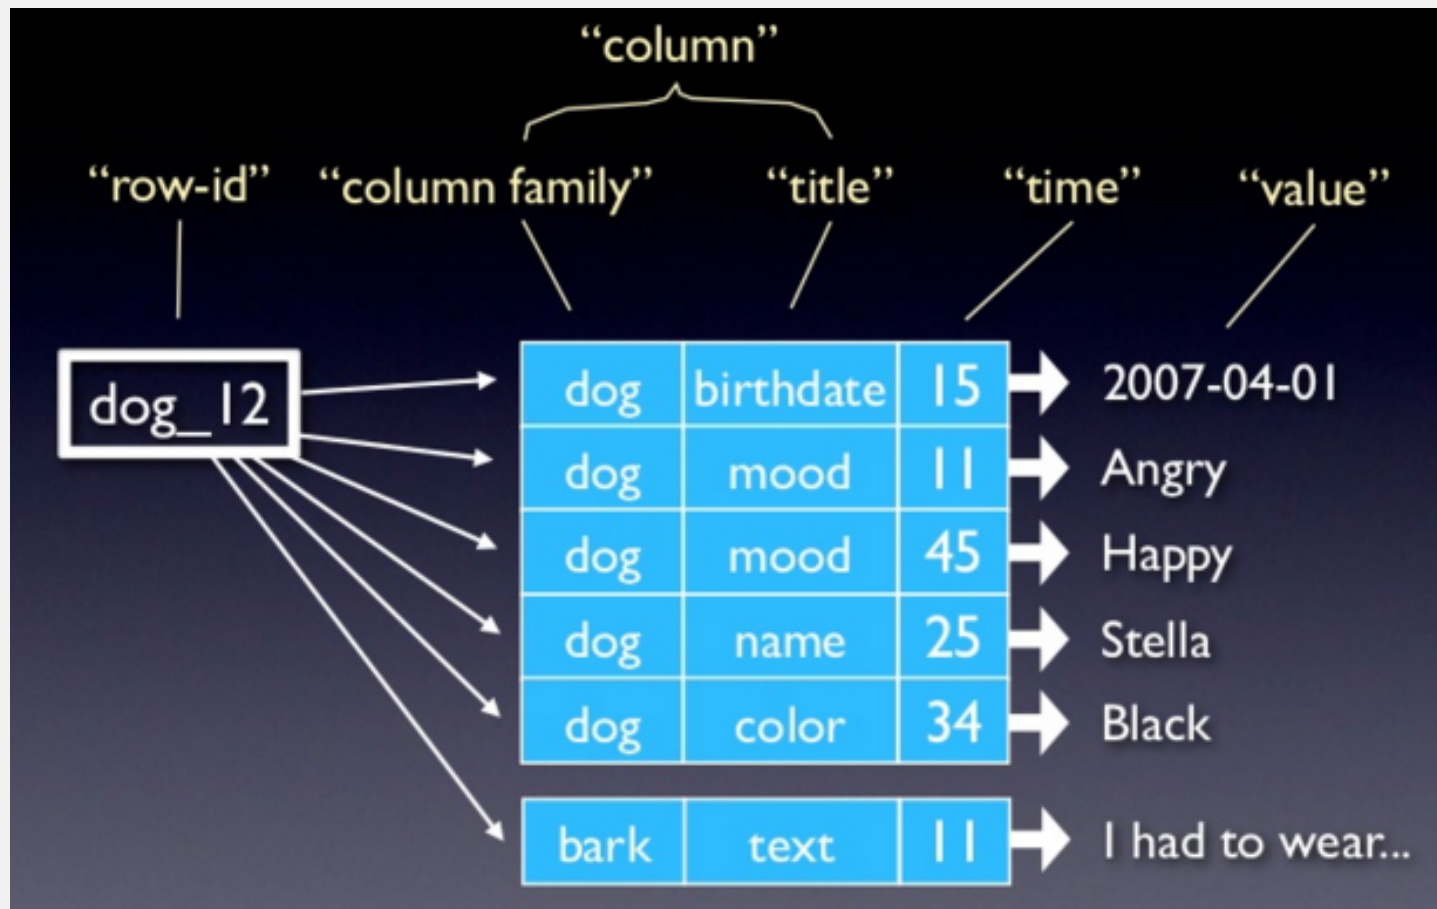
- Tables similarly to RDBMS, but handle semi-structured data
- Data model:
 - Collection of Column Families
 - Column family = (key, value) where value = set of **related** columns (standard, super)
 - indexed by *row key*, *column key* and *timestamp*
- Allow key-value pairs to be stored (and retrieved on key) as a distributed hash table
- Properties: partitioning (horizontally and/or vertically), high availability, completely transparent to application

Column-based



- One column family can have variable numbers of columns
- Cells within a column family are sorted “physically”
- Very sparse, most cells have null values
- **Comparison: RDBMS vs column-based NoSQL**
 - Query on multiple tables
 - **RDBMS:** must fetch data from several places on disk and glue together
 - **Column-based NoSQL:** only fetch column families of those columns that are required by a query
 - all columns in a column family are stored together on the disk
 - multiple rows can be retrieved in one read operation → data locality

Column-based



Column-based


- RDBMS store an entire row together (on disk or at a server)
- Column-oriented systems typically store a column together (or a group of columns).
 - Entries within a column are indexed and easy to locate, given a key (and vice-versa)
- Why useful?
 - Range searches within a column are fast since you do not need to fetch the entire database
 - E.g., Get all the `blog_ids` from the `blog` table that were updated within the past month
 - Search in the `last_updated` column, fetch corresponding `blog_id` column
 - Do not need to fetch the other columns

Column-based

Name	Producer	Data model	Querying
BigTable	Google	set of couples (key, {value})	selection (by combination of row, column, and time stamp ranges)
HBase	Apache	groups of columns (a BigTable clone)	JRUBY IRB-based shell (similar to SQL)
Hypertable	Hypertable	like BigTable	HQL (Hypertext Query Language)
CASSANDRA	Apache (originally Facebook)	columns, groups of columns corresponding to a key (supercolumns)	simple selections on key, range queries, column or columns ranges
PNUTS	Yahoo	(hashed or ordered) tables, typed arrays, flexible schema	selection and projection from a single table (retrieve an arbitrary single record by primary key, range queries, complex predicates, ordering, top-k)

Graph-based

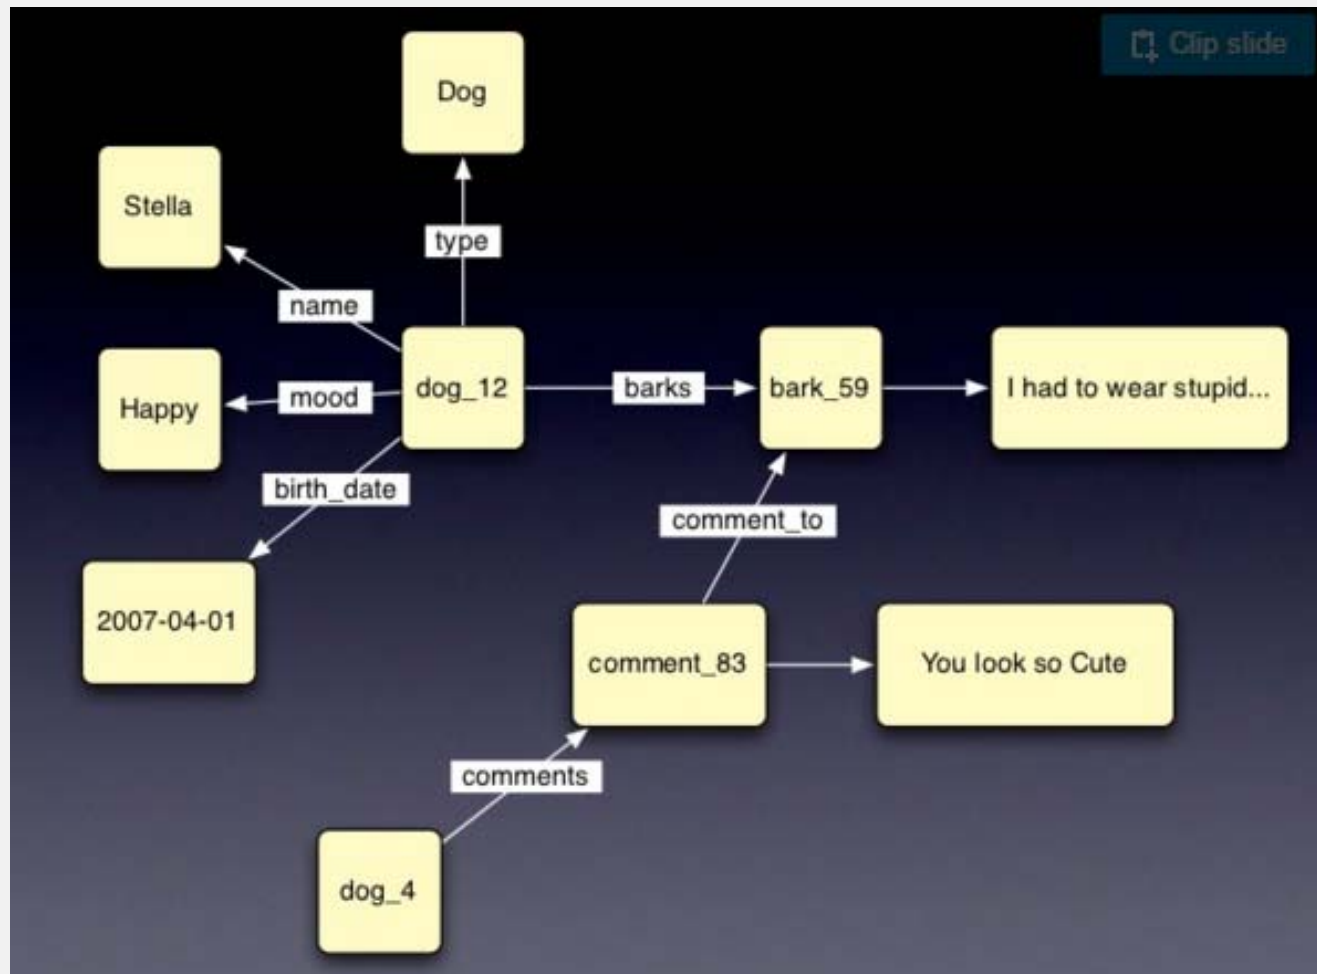
- Focus on modeling the structure of data (*interconnectivity*)
- Scales to the complexity of data
- Inspired by mathematical Graph Theory ($G=(E,V)$)
- Data model:
 - (Property Graph) nodes and edges
 - Nodes may have properties (including ID)
 - Edges may have labels or roles
 - Key-value pairs on both



Graph-based

- Good for applications where you need to traverse relationships to look for patterns such as social networks, fraud detection, and recommendation engines
- Interfaces and query languages vary
- *Single-step vs path expressions vs full recursion*
- Example:
 - Neo4j, FlockDB, Pregel, InfoGrid ...

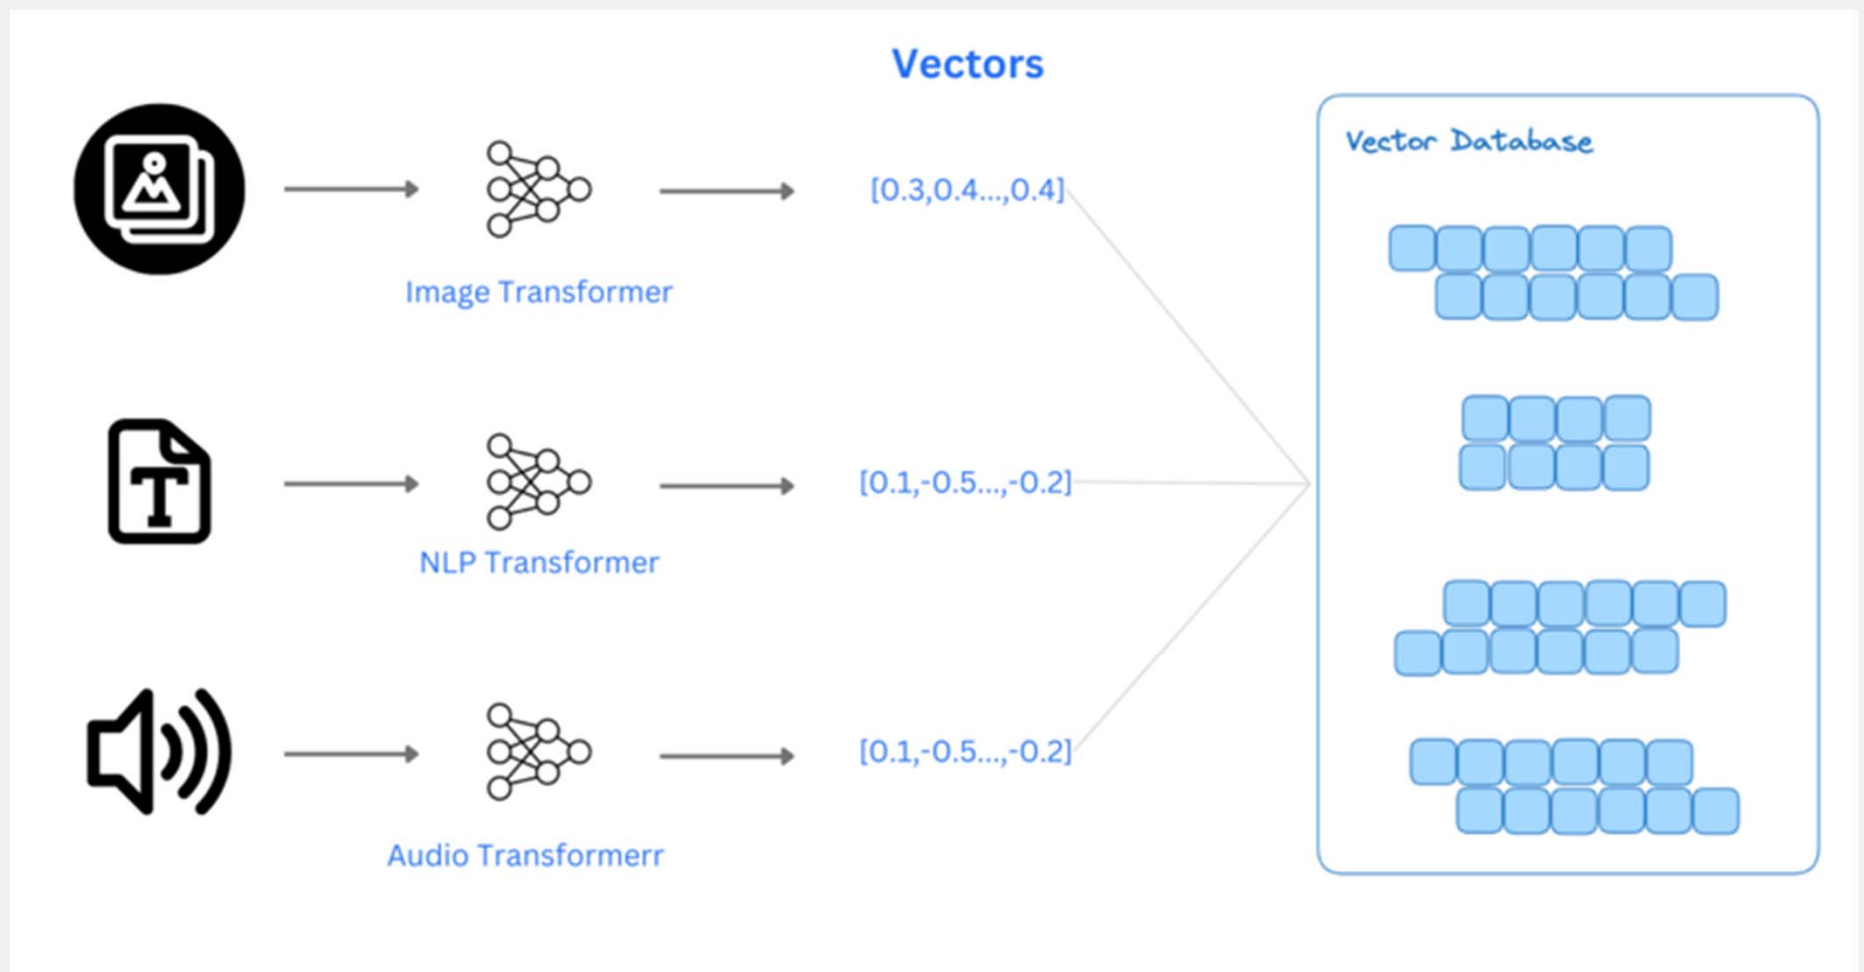
Graph-based



Vector-based

- Vector databases become popular because of large language models (LLMs)
- In LLMs, documents are represented as vectors called embeddings (e.g, Word2Vec, GloVe, transformers)
- Offer optimized storage, indexing, and querying capabilities for embeddings
- Example:
 - Pinecone, Milvus, Chroma, Weaviate, Vespa, ...

Vector-based



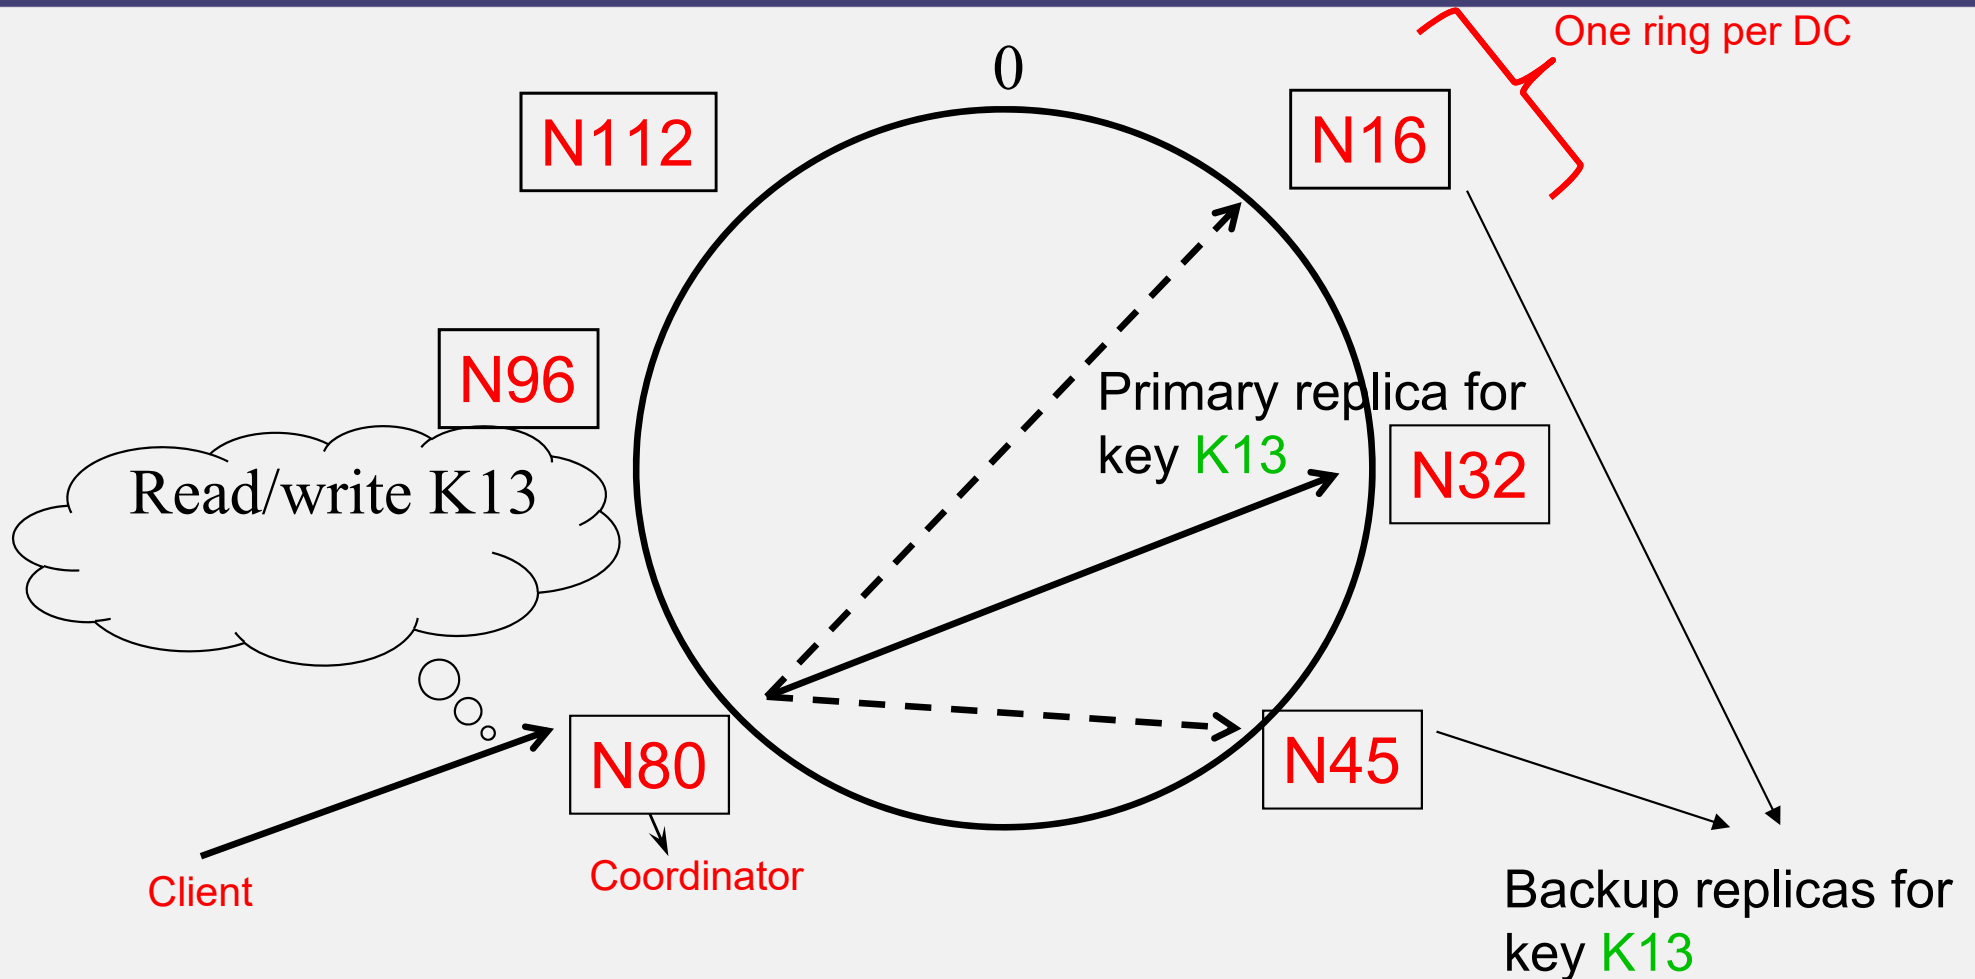
Example: Cassandra

- A distributed column-oriented key-value store
- Intended to run in a datacenter (and also across DCs)
- Originally designed at Facebook
- Open-sourced later, today an Apache project
- Some of the companies that use Cassandra in their production clusters
 - IBM, Adobe, HP, eBay, Ericsson, Symantec
 - Twitter, Spotify
 - PBS Kids
 - Netflix: uses Cassandra to keep track of your current position in the video you are watching

Let's go Inside Cassandra: Key -> Server Mapping

- How do you decide which server(s) a key-value resides on?

Key -> Server Mapping in Cassandra



Cassandra uses a Ring-based distributed hash table

Writes

- Need to be lock-free and fast (no reads or disk seeks)
- Client sends write to one coordinator node in Cassandra cluster
 - Coordinator may be per-key, or per-client, or per-query
 - Per-key Coordinator ensures writes for the key are serialized
- Coordinator uses Partitioner to send query to all replica nodes responsible for key
- When X replicas respond, coordinator returns an acknowledgement to the client
 - X? We'll see later.

Deletes

Delete: don't delete item right away

- Add a **tombstone** to the log
- Eventually, when compaction encounters tombstone it will delete item

Reads

Read: Similar to writes, except

- Coordinator can contact X replicas (e.g., in same rack)
 - Coordinator sends read to replicas that have responded quickest in past
 - When X replicas respond, coordinator returns the latest-timestamped value from among those X
 - (X? We'll see later.)
- Coordinator also fetches value from other replicas
 - Checks consistency in the background, initiating a **read repair** if any two values are different
 - This mechanism seeks to eventually bring all replicas up to date
- A row may be split across multiple tables
 - => reads need to touch multiple tables
 - => reads slower than writes (but still fast)

Membership

- Any server in cluster could be the coordinator
- So every server needs to maintain a list of all the other servers that are currently in the cluster
- List needs to be updated automatically as servers join, leave, and fail

Cassandra vs. RDBMS

Some statistics about Facebook Search (using Cassandra)

- On > 50 GB data
- MySQL
 - Writes 300 ms avg
 - Reads 350 ms avg
- Cassandra
 - Writes 0.12 ms avg
 - Reads 15 ms avg
- Orders of magnitude faster
- What's the catch? What did we lose?

Mystery of “X”: CAP Theorem

- Proposed by Eric Brewer (Berkeley)
- In a distributed system you can satisfy at most 2 out of the 3 guarantees:
 1. **Consistency**: all nodes see same data at any time, or reads return latest written value by any client
 2. **Availability**: the system allows operations all the time, and operations return quickly
 3. **Partition-tolerance**: the system continues to work in spite of network partitions (i.e., network failure)

Why is Consistency Important?

- Consistency = all nodes see same data at any time, or reads return latest written value by any client.
 - Note: this is different from the consistency in ACID
- When you access your bank or investment account via multiple clients (laptop, workstation, phone, tablet), you want the updates done from one client to be visible to other clients.
- When thousands of customers are looking to book a flight, all updates from any client should be accessible by other clients.

Why is Availability Important?

- Availability = Reads/writes complete reliably and quickly.
- Measurements have shown that a 500 ms increase in latency for operations at Amazon.com or at Google.com can cause a 20% drop in revenue.
- At Amazon, each added millisecond of latency implies a \$6M yearly loss.
- SLAs (Service Level Agreements) written by providers predominantly deal with latencies faced by clients.

Why is Partition-Tolerance Important?

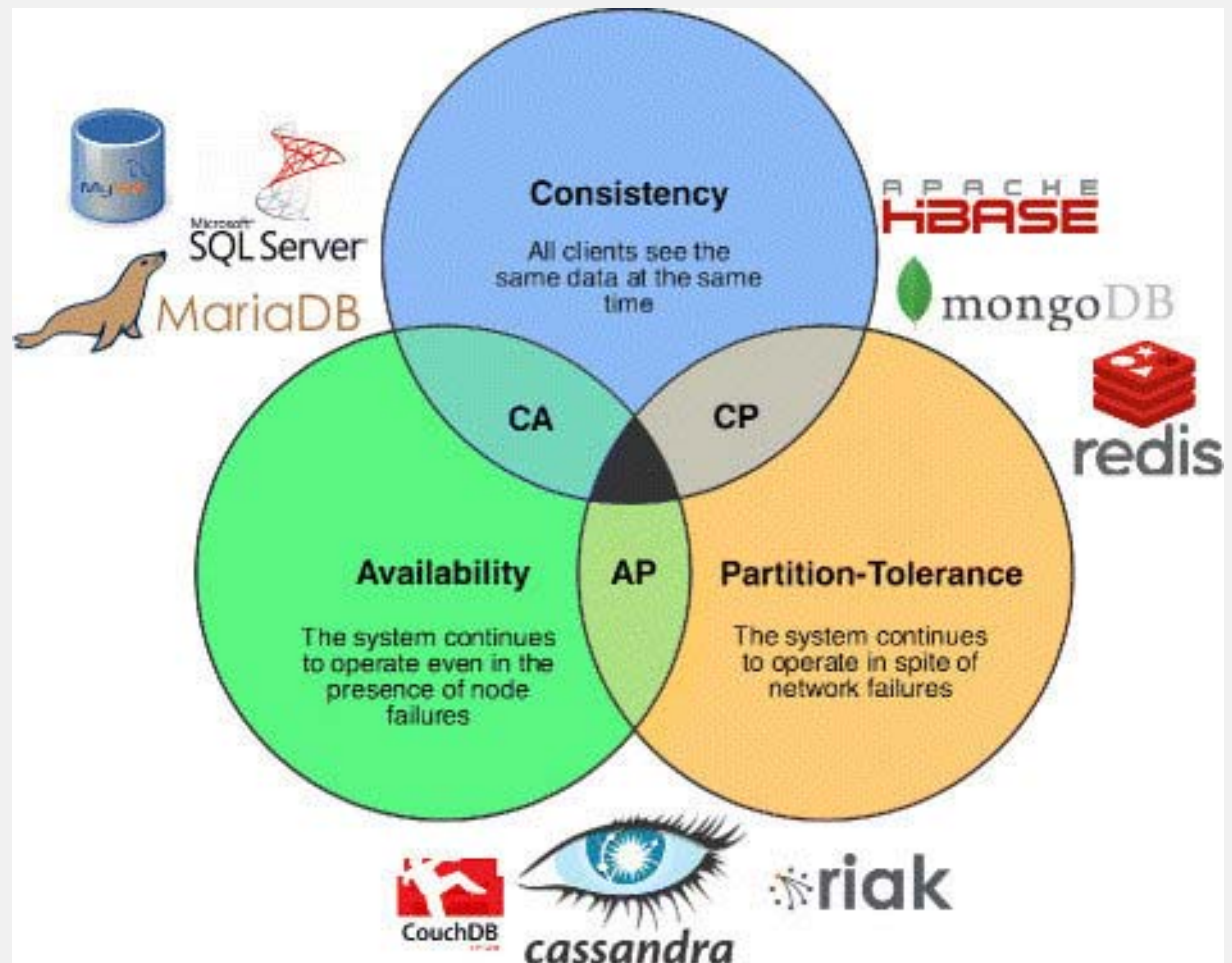
- Partitions can happen across datacenters when the Internet gets disconnected
 - Internet router outages
 - Under-sea cables cut
 - DNS not working
- Partitions can also occur within a datacenter, e.g., a rack switch outage
- Still desire system to continue functioning normally under this scenario

CAP Theorem Fallout

- Since partition-tolerance is essential in today's cloud computing systems, CAP theorem implies that a system has to choose between consistency and availability
- Cassandra
 - Eventual (weak) consistency, Availability, Partition-tolerance
- Traditional RDBMSs
 - Strong consistency over availability under a partition

CAP Tradeoff

- Starting point for NoSQL Revolution
- A distributed storage system can achieve **at most two of C, A, and P**.
- When partition-tolerance is important, you have to choose between consistency and availability



Eventual Consistency

- If all writes stop (to a key), then all its values (replicas) will converge eventually.
- If writes continue, then system always tries to keep converging.
 - Moving “wave” of updated values lagging behind the latest values sent by clients, but always trying to catch up.
- May still return stale values to clients (e.g., if many back-to-back writes).
- But works well when there a few periods of low writes—system converges quickly.


Back to Cassandra: Mystery of X

- Cassandra has **consistency levels**
- Client is allowed to choose a consistency level for each operation (read/write)
 - ANY: any server (may not be replica)
 - Fastest: coordinator caches write and replies quickly to client
 - ALL: all replicas
 - Ensures strong consistency, but slowest
 - ONE: at least one replica
 - Faster than ALL, but cannot tolerate a failure (e.g., a write operation is only confirmed by one replica)
 - QUORUM: quorum across all replicas in all datacenters (DCs)


Quorums

In a nutshell:

- Quorum = majority
 - $> 50\%$
- Then any two quorums intersect
 - e.g., Client 1 does a write in red quorum
 - Then client 2 does a read in blue quorum
 - At least one server in blue quorum returns latest write
- Quorum is faster than ALL, but still ensures strong consistency



Consistency Spectrum



Spectrum Ends: Eventual Consistency

- Cassandra offers **Eventual Consistency**
 - If writes to a key stop, all replicas of key will converge

