



MongoDB

MSBA7024 / MACC7020 Database Design and Management



What is MongoDB?

- Developed by 10gen
 - Founded in 2007
- Document-oriented, NoSQL database
 - Hash-based, *schema-less database*
 - No Data Definition Language
 - In practice, this means you can store hashes with any keys and values that you choose
 - Keys are a basic data type but in reality stored as strings
 - Document Identifiers (`_id`) will be created for each document, field name reserved by system
 - Application tracks the schema and mapping
 - Uses BSON format
- Supports APIs (drivers) in many computer languages
 - JavaScript, Python, Ruby, Perl, Java, Java Scala, C#, C++, Haskell, Erlang



Functionality of MongoDB

- Dynamic schema
 - No DDL
- Document-based database
- Secondary indexes
- Query language via an API
- Atomic writes and fully-consistent reads
 - If system configured that way
- Master-slave replication with automated failover (replica sets)
- Built-in horizontal scaling via automated range-based partitioning of data (sharding)
- No joins nor transactions


Choices Made for Design of MongoDB

- Scale horizontally over commodity hardware
 - Lots of relatively inexpensive servers
- Keep the functionality that works well in RDBMSs
 - Ad hoc queries
 - Fully featured indexes
 - Secondary indexes
- What **doesn't** distribute well in RDBMS?
 - Long running multi-row transactions
 - Joins

MongoDB: CAP approach

Focus on **Consistency** and **Partition tolerance**

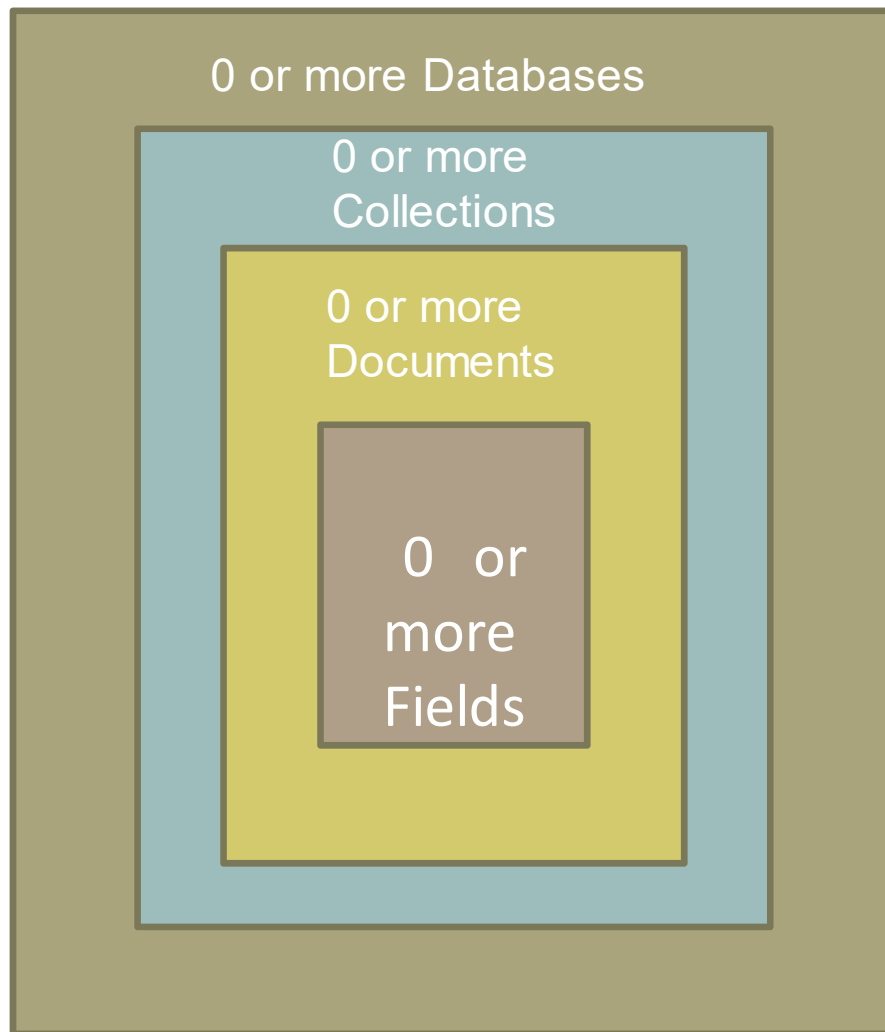
- **Consistency**
 - all replicas contain the same version of the data
- **Availability**
 - system remains operational on failing nodes
- **Partition tolerance**
 - multiple entry points
 - system remains operational on system split



CAP Theorem:
it is **impossible** to satisfy all three
at the same time

MongoDB: Hierarchical Objects

- A MongoDB instance may have zero or more 'databases'
- A database may have zero or more 'collections'
- A collection may have zero or more 'documents'
- A document may have one or more 'fields'
- MongoDB 'Indexes' function much like their RDBMS counterparts



RDBMS Concepts to NoSQL

RDBMS		MongoDB
Database	➔	Database
Table, View		Collection
Row		Document (BSON)
Column		Field
Index		Index
Join		Embedded Document
Foreign Key		Reference
Partition		Shard

Collection is not strict about what it stores

Schema-less

Hierarchy is evident in the design

Embedded Document

MongoDB Processes and Configuration

- mongod – Database instance
- mongos – Sharding processes
 - Analogous to a database router
 - Processes all requests
 - Decides how many and which *mongods* should receive the query
 - *mongos* collates the results, and sends it back to the client

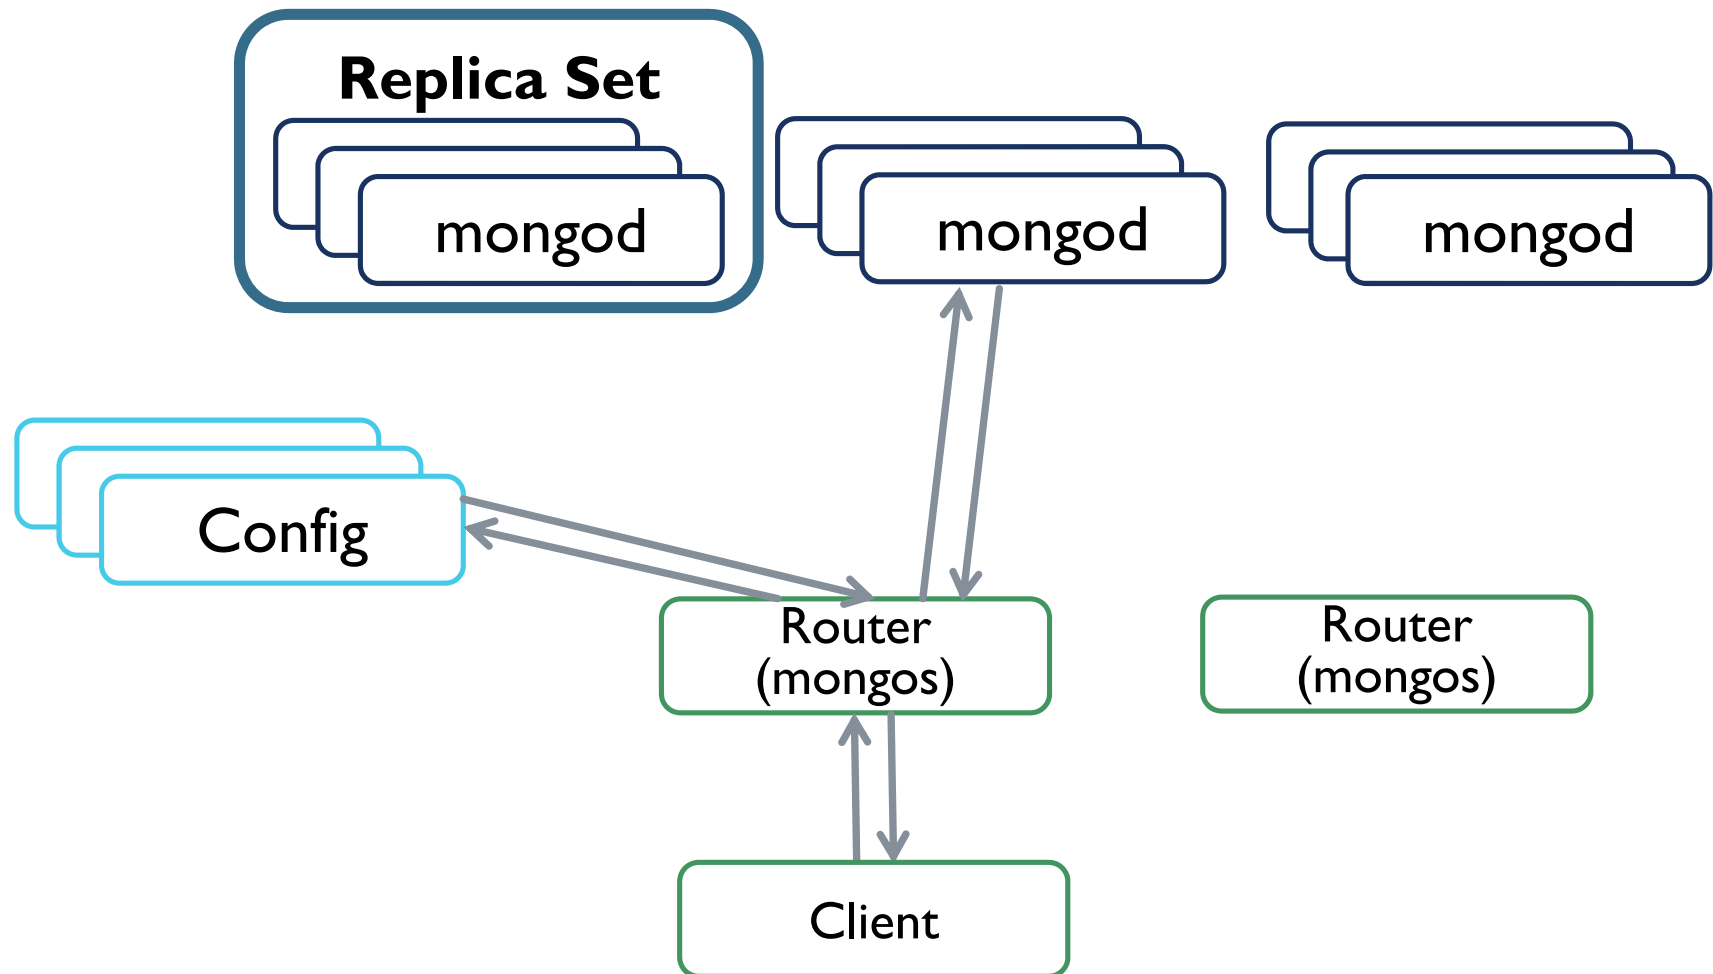
MongoDB Processes and Configuration

- Data split into chunks, based on shard key (~ primary key)
- Either use hash or range-partitioning
- Shard: collection of chunks
- Shard assigned to a replica set

MongoDB Processes and Configuration

- Replica set consists of multiple mongod servers (typically 3 mongods)
- Replica set members are mirrors of each other
 - One is primary
 - Others are secondary
- You can have one *mongos* for the whole system no matter how many mongods you have
- Or you can have one local *mongos* for every client if you want to minimize network latency.

MongoDB Processes and Configuration



MongoDB Processes and Configuration

- MongoCompass
 - A graphical client to interact with the MongoDB server
- Mongo – an interactive shell (also a client)
 - Fully functional JavaScript environment for use with a MongoDB

Schema Free

- MongoDB does not need any pre-defined data schema
- Every document in a collection could have different data
 - Addresses NULL data fields

```
{name: "will",  
  eyes: "blue",  
  birthplace: "NY",  
  aliases: ["bill", "la ciacco"],  
  loc: [32.7, 63.4],  
  boss: "ben"}
```

```
{name: "jeff",  
  eyes: "blue",  
  loc: [40.7, 73.4],  
  boss: "ben"}
```

```
{name: "brendan",  
  aliases: ["el diablo"]}
```

```
{name: "ben",  
  hat: "yes"}
```

```
{name: "matt",  
  pizza: "DiGiorno",  
  height: 72,  
  loc: [44.6, 71.3]}
```



BSON Format

- Binary-encoded serialization of JSON-like documents
- Zero or more key/value pairs are stored as a single entity
- Each entry consists of a field name, a data type, and a value
- Large elements in a BSON document are prefixed with a length field to facilitate scanning

JSON Format

- Data is in name / valuepairs
- A name/value pair consists of a field name followed by a colon, followed by a value:
 - Example: "name": "R2-D2"
- Data is separated by commas
 - Example: "name": "R2-D2", race : "Droid"
- Curly braces hold objects
 - Example: {"name": "R2-D2", race : "Droid", affiliation: "rebels"}
- An array is stored in brackets []
 - Example: [{"name": "R2-D2", race : "Droid", affiliation: "rebels"}, {"name": "Yoda", affiliation: "rebels"}]

JSON Format

- Example of a document

```
{  
  name: "travis",  
  salary: 30000,  
  designation: "Computer Scientist",  
  teams: [ "front-end", "database" ]  
}
```


Index Functionality

- B+ tree indexes
- An index is automatically created on the `_id` field (the primary key)
- Users can create other indexes to improve query performance or to enforce Unique values for a particular field
- Supports single field index as well as Compound index
 - Like SQL order of the fields in a compound index matters
 - If you index a field that holds an array value, MongoDB creates separate index entries for *every* element of the array
- Sparse property of an index ensures that the index only contain entries for documents that have the indexed field (so ignore records that do not have the field defined)
- If an index is both unique and sparse – then the system will reject records that have a duplicate key value but allow records that do not have the indexed field defined

CRUD Operations

- Create
 - `db.collection_name.insert(<document>)`
 - `db.collection_name.save(<document>)`
- Read
 - `db.collection_name.find(<query>, <projection>)`
 - `db.collection_name.findOne(<query>, <projection>)`
- Update
 - `db.collection_name.update(<query>, <update>, <options>)`
- Delete
 - `db.collection_name.deleteOne(<query>)`
 - `db.collection_name.deleteMany(<query>)`

`db.collection` **specifies** the collection or the 'table' for the operation.

Create Operations

db.collection specifies the collection or the 'table' to store the document

- db.collection_name.insert(<document>)
 - Omit the _id field to have MongoDB generate a unique key
 - Example
 - db.**parts**.insert({type: "screwdriver", quantity: 15 })
 - db.**parts**.insert({_id: 10, type: "hammer", quantity: 1 })
- db.collection_name.update(<query>, <update>, { upsert: true })
 - Will update 1 or more records in a collection satisfying query
- db.collection_name.save(<document>)
 - Updates an existing record or creates a new record

Create Operations

Insert a row entry for new employee Sally

```
db.employee.insert({  
    name : "sally",  
    salary : 15000,  
    designation : "MTS",  
    teams : [ "cluster-management" ]  
})
```

Read Operations

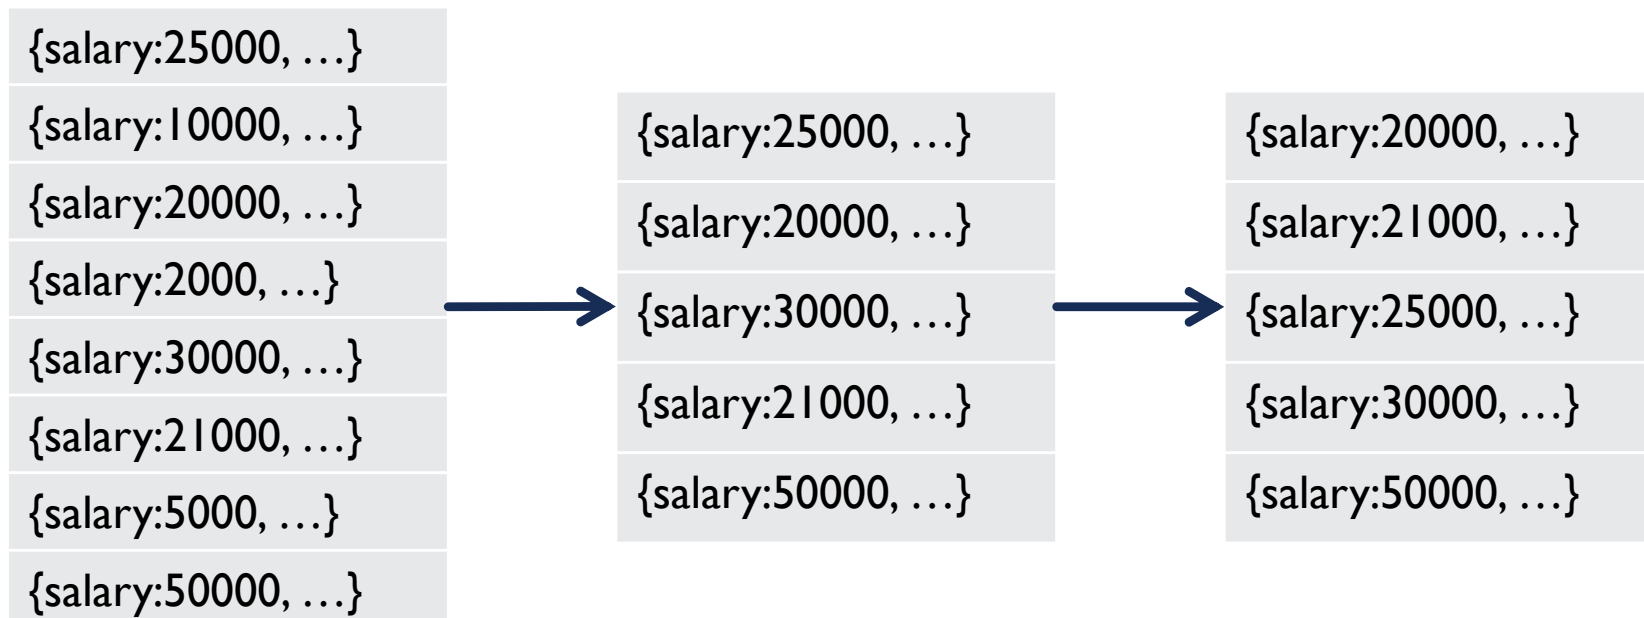
- `db.collection_name.find(<query>, <projection>)`
 - Provides functionality similar to the SELECT command
 - `<query>` where condition, `<projection>` fields in result set
 - Example
 - `db.parts.find({type:"hammer"}).limit(5)`
 - Can use a cursor to handle a result set
 - `var PartsCursor = db.parts.find({type:"hammer"}).limit(5)`
 - Can modify the query to impose limits, skips, and sortorders.
 - Can specify to return the 'top' number of records from the result set
- `db.collection_name.findOne(<query>, <projection>)`

Read Operations

Query all employee names with salary greater than 18000 sorted in ascending order

```
db.employee.find({salary:{$gt:18000}, {name:1}}).sort({salary:1})
```

Collection Condition Projection Modifier



Query Operators

Name	Description
\$eq	Matches value that are equal to a specified value
\$gt, \$gte	Matches values that are greater than (or equal to a specified value
\$lt, \$lte	Matches values less than or (equal to) a specified value
\$ne	Matches values that are not equal to a specified value
\$in	Matches any of the values specified in an array
\$nin	Matches none of the values specified in an array
\$or	Joins query clauses with a logical OR returns all
\$and	Join query clauses with a logical AND
\$not	Inverts the effect of a query expression
\$nor	Join query clauses with a logical NOR
\$exists	Matches documents that have a specified field

Query Operators

- `db.employee.find()`
- `db.employee.find({name: "Sally"})`
- `var cursor = db.employee.find({salary: {$in: [5000, 2000] } })`
 - Use `next()` to access the rest of the records

Query Operators

- `db.employee.find({name: "Steve", salary: { $lt: 3000 } })`
- `db.employee.find({ $or: [{ name: "Bill" }, { salary: { $gt: 9000 } }] })`
- Find records of all managers who earn more than 5000
`db.employee.find({designation:"Manager", salary: { $gt: 5000 } })`

Update Operations

- `db.collection_name.update(<query>, <update>, <options>)`
 - Will update 1 or more records in a collection satisfying query
- `db.collection_name.findAndModify(<query>, <sort>, <update>, <new>, <fields>,<upsert>)`
 - Modify existing record(s) – retrieve old or new version of the record

Update Operations

All employees with salary greater than 18000 get a designation of Manager

<i>Update Criteria</i>	<code>db.employee.update(</code>
<i>Update Action</i>	<code> {salary:{\$gt:18000}},</code>
<i>Update Option</i>	<code> {\$set: {designation: "Manager"}},</code>
	<code> {multi: true}</code>
	<code>)</code>

Multi-option allows multiple document update.

Update Operations

- Increment salary of all managers by 1000
 - `db.employee.update({ designation : "Manager" }, { $inc : { salary : 1000 } })` -> update one only
 - `db.employee.update({ designation : "Manager" }, { $inc : { salary : 1000 } } , { multi: true })`
- Increment salary of all managers working in cluster-management group by 5000
 - `db.employee.update({ designation : "Manager", teams: "cluster-management" }, { $inc : { salary : 5000 } } , { multi: true })`

Delete Operations

- `db.collection_name.deleteOne(<query>)`
 - Delete the first record from a collection or matching a criterion
 - `db.parts.deleteOne({ type: /^h/ })` – remove all parts starting with h
- `db.collection_name.deleteMany(<query>)`
 - Delete all records from a collection or matching a criterion

Delete Operations

Remove all employees who earn less than 10000

Remove Criteria

```
db.employee.deleteMany(  
  {salary:{$lt:10000}},  
)
```

Aggregated Functionality

Aggregation framework provides SQL-like aggregation functionality

- Pipeline documents from a collection pass through an aggregation pipeline, which transforms these objects as they pass through
- Expressions produce output documents based on calculations performed on input documents
- Example: `db.parts.aggregate ({$group : {_id: "$type", totalquantity : { $sum: "$quantity"} } })`

Document Structures

- Two main document structures in MongoDB
 - Embedded Data
 - References

Document Structures

- Embedded Data: capture relationships between data by storing related data in a single document structure.

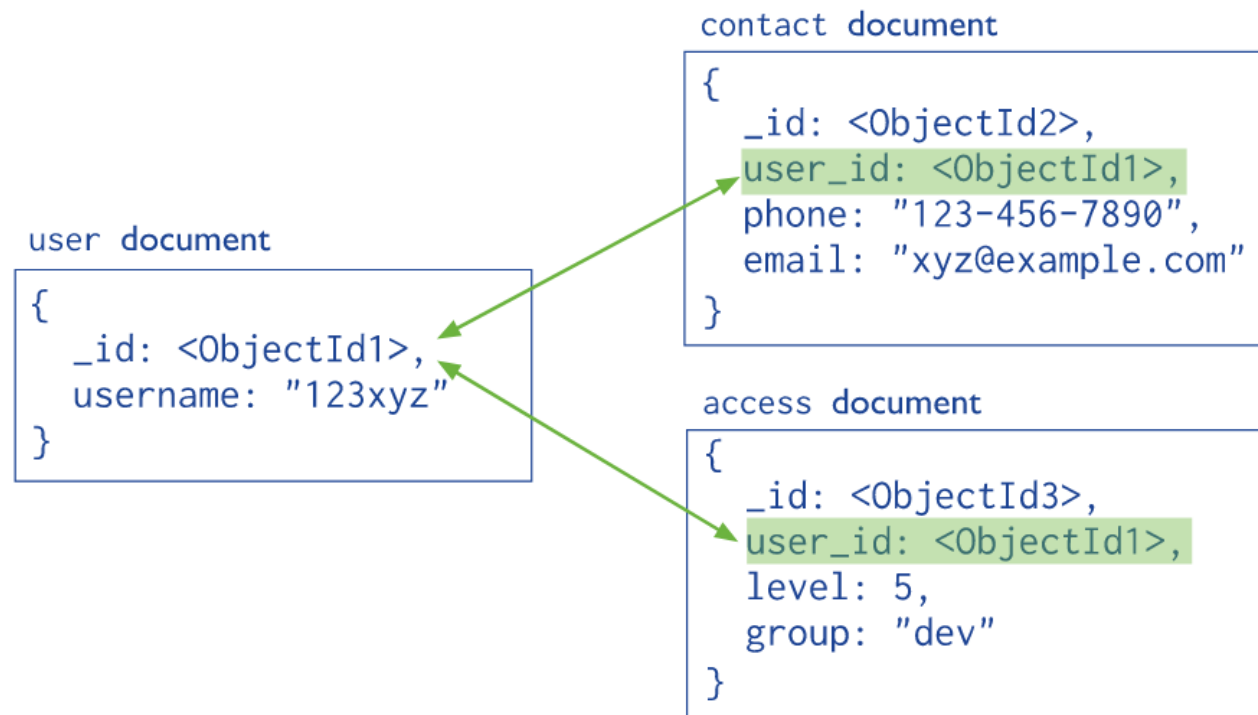
```
{
  _id: <ObjectId>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

Embedded sub-document

Embedded sub-document

Document Structures

- References: store the relationships between data by including links or references from one document to another.
- Normalized data models



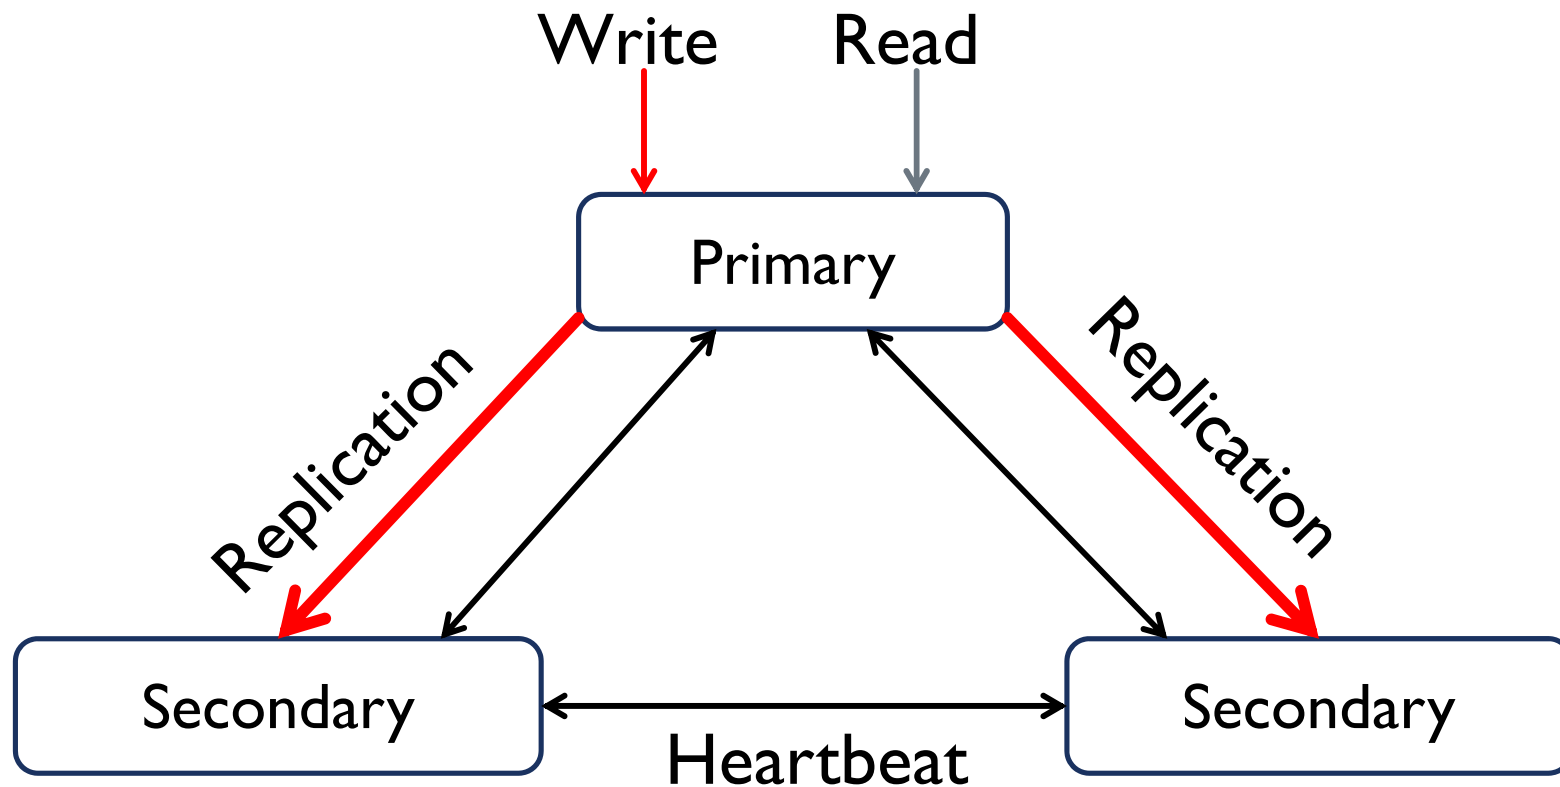
Document Structures

- Embedded data
 - Fast retrieval – retrieve all related data in one read operation
 - Normally used for one-to-one and one-to-many relationships between entities
 - May have duplicated data
- References
 - Slow retrieval – need multiple read operations
 - Can model more complex many-to-many relationship or large hierarchical data

Document Structures

- Embedded document pattern can lead to large documents
 - e.g., a product with several thousand reviews
- Subset Pattern
 - Instead of storing all of the reviews with the product, we split the collection into two collections:
 - The product collection stores information on each product, including the product's ten most recent reviews
 - The review collection stores all reviews. Each review contains a reference to the product for which it was written.
 - Fast to see the product page with top 10 reviews
 - If a user wants to see additional reviews, the application makes a call to the review collection.
- Need to consider:
 - Data duplication
 - Make sure the subset is the desired/correct subset

Replication of Data



(the nodes send pings to each other every two seconds to identify the current stats)

Replication of Data

- Ensures redundancy, backup, and automatic failover
- Replication occurs through groups of servers known as replica sets
 - Primary set – set of servers that client tasks directly updates to
 - Secondary set – set of servers used for duplication of data
 - At most can have 12 replica sets
 - Many different properties can be associated with a secondary set i.e. secondary-only, hidden delayed, arbiters, non-voting
 - If the primary set fails, the secondary sets 'vote' to elect the new primary set

Read Preference

- Determine where to route read operation
- Default is primary. Some other options are
 - primary-preferred
 - secondary
 - nearest
- Helps reduce latency, improve throughput
- Reads from secondary may fetch stale data

Write Concern

- Determines the guarantee that MongoDB provides on the success of a write operation
- Default is *acknowledged* (primary returns answer immediately).
- Other options are
 - journaled (typically at primary)
 - replica-acknowledged (quorum with a value of W), etc.
- Weaker write concern implies faster write time

Consistency of Data

- All read operations issued to the primary of a replica set are consistent with the last write operation
 - Reads to a primary have **strict consistency**
 - Reads reflect the latest changes to the data
 - Reads to a secondary have **eventual consistency**
 - Updates propagate gradually
- If clients permit reads from secondary sets – then client may read a previous state of the database
- Failure occurs before the secondary nodes are updated
 - System identifies when a rollback needs to occur
 - Users are responsible for manually applying rollback changes