


Kernel SVM

Hyperplane

- In a p -dim feature space, a hyperplane has the form

$$\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p = 0$$



E.g. 2-dim space


$$\beta_0 + \beta_1 x_1 + \beta_2 x_2 = 0$$

What is $(\beta_0, \beta_1, \beta_2)$?

Hyperplane

- Distance to hyperplane (when $\beta_0 = 0$) E.g. $(x_1, x_2) = (2, 0)$


$$\vec{x} \cdot \vec{\beta}$$



- Note: Distance can be negative

Hyperplane


- When $\beta_0 \neq 0$: shift the hyperplane along the direction of normal vector by $|\beta_0|$



Hyperplane

- Distance to hyperplane (when $\beta_0 \neq 0$)


$$\vec{x} \cdot \vec{\beta} + \beta_0$$



- Note: Again, distance can be negative

Hyperplane

- Margin M : two lines parallel to the hyperplane
- Classification via [Hyperplane with Margin]



$$\forall i \text{ s.t. } y_i = +1, \quad \vec{x}^{(i)} \cdot \vec{\beta} + \beta_0 \geq M$$

$$\forall i \text{ s.t. } y_i = -1, \quad \vec{x}^{(i)} \cdot \vec{\beta} + \beta_0 \leq -M$$

Data Separable

- Solution to the optimization problem: $\max_{\beta_0, \beta_1, \dots, \beta_p, M} M$

$$\text{s.t.} \quad \sum_{j=1}^p \beta_j^2 = 1$$

$$y_i(\beta_0 + \beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_p x_p^{(i)}) \geq M$$


Note: Given $(\beta_0, \vec{\beta})$, hyperplane is fixed \rightarrow Max M can be easily obtained

- Expand the two margin lines until one touches a data point

Thus, just search through the space of all $(\beta_0, \vec{\beta}) \rightarrow$ Find the optimal M

Data Separable

- Support vectors:
 - With respect to the optimal hyperplane, points that lie on the margin lines



How many support vectors are there?

Data Non-Separable

- Solution to the optimization problem $\max_{\beta_0, \beta_1, \dots, \beta_p, \xi_i, M} M$

$$\text{s.t.} \quad \sum_{j=1}^p \beta_j^2 = 1$$

$$y_i(\beta_0 + \beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_p x_p^{(i)}) \geq M(1 - \xi_i)$$

$$\xi_i \geq 0, \sum_{i=1}^N \xi_i \leq C$$

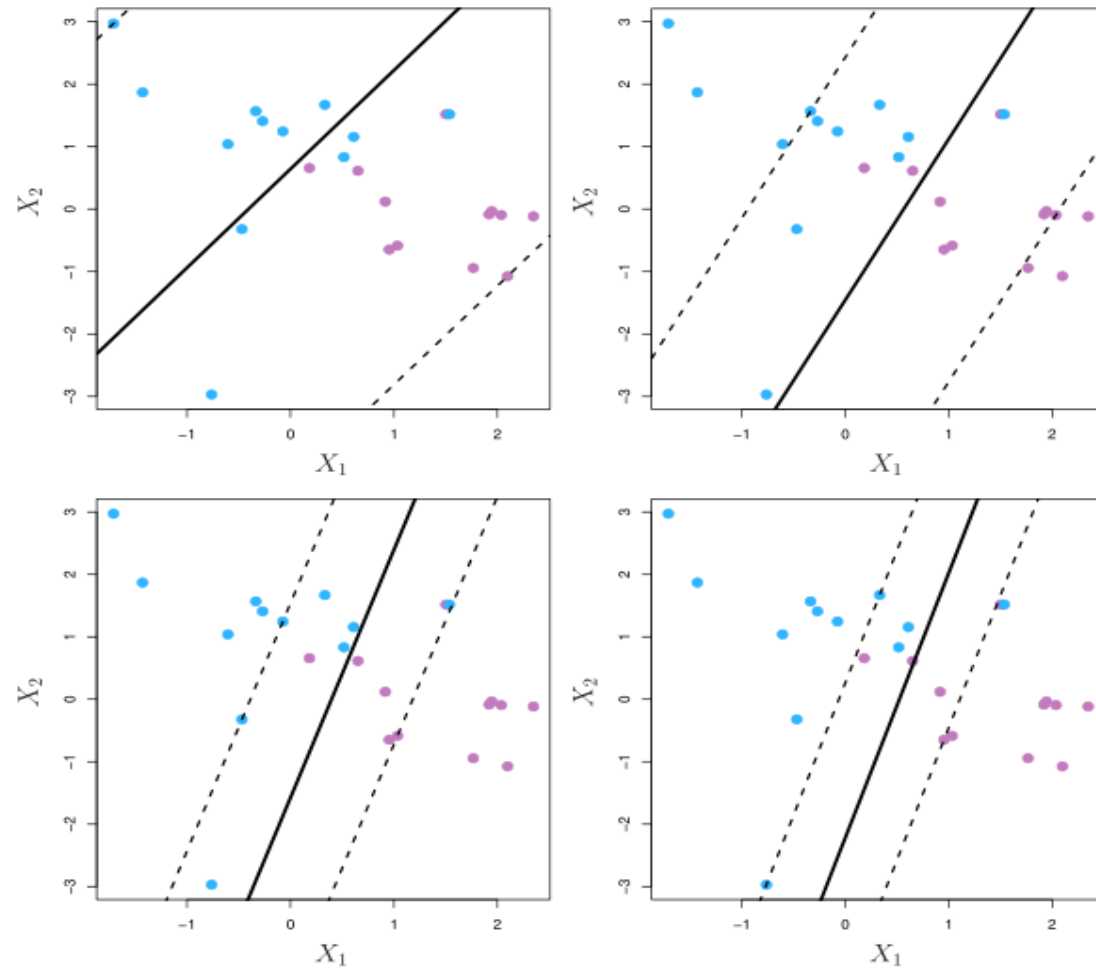
where ξ_i 's are the slack variables, C is a nonnegative tuning parameter

- Input: [N by p] X (training feature matrix), [N by 1] Y (training label vector)
- Optimization problem can be solved if N and p are both NOT too large

The role of C

- C determines the number and severity of the violations to the margin (and to the hyperplane) that we will tolerate.

The role of C : Bias-variance Tradeoff




- Which figure has the largest C ?

Nonlinear SVM & Kernels

Non-linear Decision Boundaries: Motivation

- Linear SVM: good approach if **decision boundary linear**
- In practice, many problems have non-linear boundaries



- Obvious nonlinear boundary:
 - Linear SVM (right figure) performs poorly here

Non-linear Decision Boundaries: Motivation

How to obtain nonlinear boundaries?

Non-linear Decision Boundaries


For instance, we can try the following

- Previous: x_1, x_2, \dots, x_p
- Now: $x_1, x_2, \dots, x_p, x_1^2, x_2^2, \dots, x_p^2$
- Optimization

$$\begin{aligned} \max_{\beta, \xi_i} M \\ \sum_{j=1}^p \beta_{j1}^2 + \sum_{j=1}^P \beta_{j2}^2 &= 1 \\ y_i(\beta_0 + \sum_{j=1}^P \beta_{j1} x_{ij} + \sum_{j=1}^P \beta_{j2} x_{ij}^2) &\geq M(1 - \xi_i) \\ \xi_i &\geq 0, \sum_{i=1}^N \xi_i \leq C \end{aligned}$$

- Why is this a nonlinear boundary?

Non-linear Decision Boundaries



Source: <https://www.youtube.com/watch?v=9NrALgHFwTo>

Non-linear Decision Boundaries

- We tried: **p features**: $x_1, x_2, \dots, x_p \rightarrow$ **$2p$ features**: $x_1, x_2, \dots, x_p, x_1^2, x_2^2, \dots, x_p^2$
- May also want to include
 - higher-order polynomial terms / interaction terms e.g. $x_1 x_2$ / other funcs
- Problem: too many predictors
 - E.g. Include up to deg- d polynomials: from p features to $O(p^d)$ features
 - Quickly become computational unmanageable
- We will show how to: enlarge feature space & computationally efficient

Kernel Trick Preliminaries

- $\phi: p\text{-dim} \rightarrow q\text{-dim} (q > p)$
 - $x^{(i)} \rightarrow \phi(x^{(i)})$, i.e. ϕ transforms $x^{(i)}$ to the enlarged feature space, $i = 1, 2, \dots, N$
 - E.g. ($p = 2, q = 3$): $x^{(i)} = (x_1^{(i)}, x_2^{(i)})$, $\phi((x_1^{(i)}, x_2^{(i)})) = (x_1^{(i)}, x_2^{(i)}, (x_1^{(i)})^2 + (x_2^{(i)})^2)$
- Kernel K : $p\text{-dim} * p\text{-dim} \rightarrow 1\text{-dim}$
 - $K(x^{(i)}, x^{(j)}) = \phi(x^{(i)}) \cdot \phi(x^{(j)})$
- ϕ is our enlarged feature space, i.e. has dim $q > p$. (In fact, it can be ∞ -dimensional)
- This seems crazy, \uparrow dim from p to ∞ (p to p^d already NOT manageable for d large)
- Let's see how we can actually do that
 - 1. Important property of SVM optimization problem
 - 2. Kernel trick

Important property of Linear SVM

- Solution to the optimization problem

$$\text{s.t.} \quad \sum_{j=1}^p \beta_j^2 = 1$$

$$y_i(\beta_0 + \beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \cdots + \beta_p x_p^{(i)}) \geq M(1 - \xi_i)$$

$$\xi_i \geq 0, \sum_{i=1}^N \xi_i \leq C$$

where ξ_i 's are the slack variables, C is a nonnegative tuning parameter

- Important Property: solution to the above involves only $x^{(i)} \cdot x^{(j)}$, $1 \leq i, j \leq N$

(This slide is intended to be empty)

Important property of Linear SVM

- For the enlarged feature space (possibly ∞ -dim), we just need $K(x^{(i)}, x^{(j)})$ for all pairs of $(x^{(i)}, x^{(j)})$
- Kernel trick: compute $K(x^{(i)}, x^{(j)}) = \phi(x^{(i)}) \cdot \phi(x^{(j)})$ without using $\phi(x^{(i)})$ or $\phi(x^{(j)})$

Examples (Optional)

- $\phi((x_1^{(i)}, x_2^{(i)})) = ((x_1^{(i)})^2, (x_2^{(i)})^2, \sqrt{2}x_1^{(i)}x_2^{(i)})$

Examples (Optional)

- $\phi((x_1^{(i)}, x_2^{(i)})) = (1, (x_1^{(i)})^2, (x_2^{(i)})^2, \sqrt{2}x_1^{(i)}, \sqrt{2}x_2^{(i)}, \sqrt{2}x_1^{(i)}x_2^{(i)})$

Examples (Optional)

- $K(x^{(i)}, x^{(j)}) = e^{-\gamma |x^{(i)} - x^{(j)}|^2}$

Kernel functions

Suppose $u, v \in \mathbb{R}_p$. Popular kernel functions include:

- **Linear kernel**: $K(u, v) = u^T v + c, c \in \mathbb{R}$ is a constant **Linear SVM**
- **Polynomial kernel**: $K(u, v) = (1 + u^T v)^d, d > 0$
- **Gaussian kernel or radial basis function kernel**: $K(u, v) = e^{-\gamma|u-v|^2}, \gamma > 0$


SVM Implementation

```
# Kernel SVM
library(e1071)
set.seed(1)
x=matrix(rnorm(200*2), ncol=2)
x[1:100,]=x[1:100,]+2
x[101:150,]=x[101:150,]-2
y=c(rep(1,150),rep(2,50))
dat=data.frame(x=x,y=as.factor(y))
plot(x, col=y)
```

**rnorm: random # from
standard normal dist., then
resize into 200 by 2 matrix**

SVM Implementation

```
# Kernel SVM  
library(e1071)  
set.seed(1)  
x=matrix(rnorm(200*2), ncol=2)  
x[1:100,]=x[1:100,]+2  
x[101:150,]=x[101:150,]-2  
y=c(rep(1,150),rep(2,50))  
dat=data.frame(x=x,y=as.factor(y))  
plot(x, col=y)
```



SVM Implementation

```
train=sample(200,100)
```

$$K(u, v) = e^{-\gamma \|u - v\|_2^2}$$


Radial basis kernel with $\gamma = 1$, $C = 1$

```
# SVM with normal C
```

```
svmfit_normalC=svm(y~., data=dat[train,], kernel="radial", gamma=1, cost=1)
```

```
plot(svmfit_normalC, dat[train,]) # plot svm, note support vector marked as crosses
```

SVM classification plot



SVM Implementation

```
train=sample(200,100)
```

$$K(u, v) = e^{-\gamma \|u-v\|_2^2}$$

Radial basis kernel with gamma = 1, C = 1

```
# SVM with normal C
svmfit_normalC=svm(y~., data=dat[train,], kernel="radial", gamma=1, cost=1)
plot(svmfit_normalC, dat[train,]) # plot svm, note support vector marked as crosses
summary(svmfit_normalC)
table(true=dat[train,"y"], pred=predict(svmfit_normalC,newdata=dat[train,]))
table(true=dat[-train,"y"], pred=predict(svmfit_normalC,newdata=dat[-train,]))
```

```
> summary(svmfit_normalC)
```

Summary of model, train & test error

Call:

```
svm(formula = y ~ ., data = dat[train, ], kernel = "radial", gamma = 1, cost = 1)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1
```

Number of Support Vectors: 31

```
( 16 15 )
```

Number of Classes: 2

Levels:

```
1 2
```

```
> table(true=dat[train,"y"], pred=predict(svmfit_normalC,newdata=dat[train,]))
```

```
      pred
true  1  2
1    69  4
2     3 24
```

```
> table(true=dat[-train,"y"], pred=predict(svmfit_normalC,newdata=dat[-train,]))
```

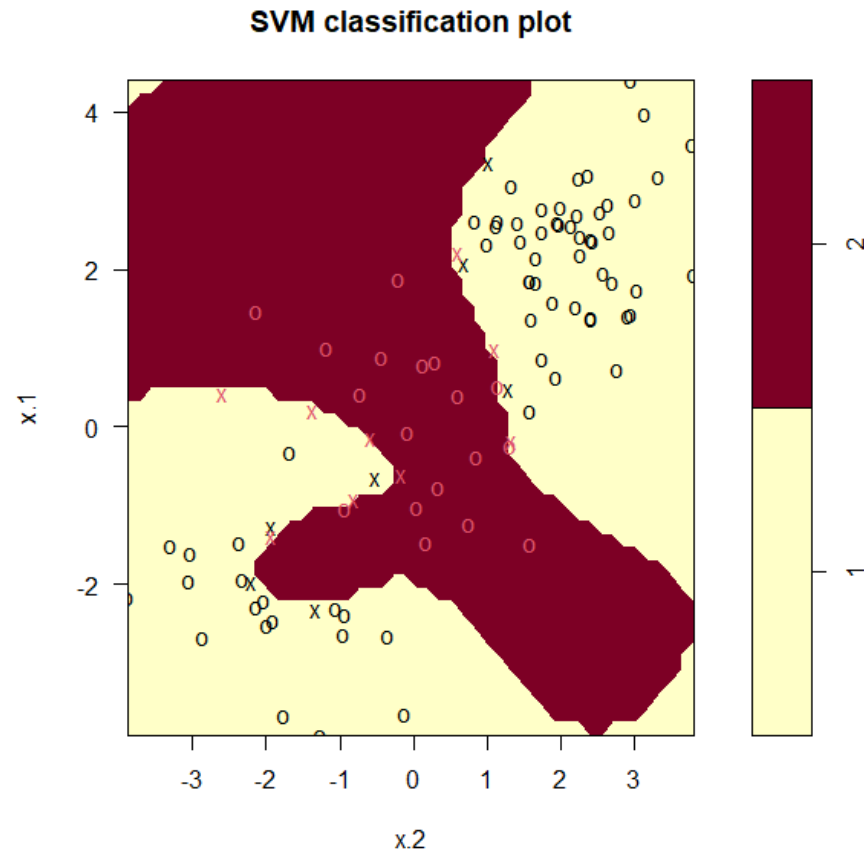
```
      pred
true  1  2
1    67 10
2     3 20
```

Train error: 7%

Test error: 13%

SVM Implementation

```
# SVM with large C (small margin)
svmfit_largeC=svm(y~., data=dat[train,], kernel="radial",gamma=1,cost=1e5)
plot(svmfit_largeC, dat[train,], svSymbol = "x")
```



SVM Implementation

```
# SVM with large C (small margin)
svmfit_largeC=svm(y~., data=dat[train,], kernel="radial",gamma=1,cost=1e5)
plot(svmfit_largeC, dat[train,], svSymbol = "x")
summary(svmfit_largeC)
table(true=dat[train,"y"], pred=predict(svmfit_largeC,newdata=dat[train,]))
table(true=dat[-train,"y"], pred=predict(svmfit_largeC,newdata=dat[-train,]))
```

```
> summary(svmfit_largeC)
```

Call:
svm(formula = y ~ ., data = dat[train,], kernel = "radial", gamma = 1, cost = 1e+05)

Parameters:
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1e+05

Number of Support Vectors: 16

(7 9)

Number of Classes: 2

Levels:
1 2

Perform perfectly on the training set, poorly on the test set (overfit)

```
> table(true=dat[train,"y"], pred=
      pred
true  1  2
      1 73  0
      2  0 27
> table(true=dat[-train,"y"], pred=
      pred
true  1  2
      1 57 20
      2  5 18
```

SVM Implementation

```
## Tune parameters
```

```
set.seed(1)
```

```
tune.out=tune(svm, y~., data=dat[train,], kernel="radial",  
              ranges=list(cost=c(0.1,1,10,100,1000), gamma=c(0.5,1,2,3,4)))
```

```
summary(tune.out)
```

Note: tune() performs 10-fold CV

```
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost	gamma
1	0.5

- best performance: 0.07

- Detailed performance results:

	cost	gamma	error	dispersion
1	1e-01	0.5	0.26	0.15776213
2	1e+00	0.5	0.07	0.08232726
3	1e+01	0.5	0.07	0.08232726
4	1e+02	0.5	0.14	0.15055453
5	1e+03	0.5	0.11	0.07378648

**CV error of best
model: 7%**

SVM Implementation

```
library(dplyr)
arrange(tune.out$performances,error)

table(true=dat[-train,"y"], pred=predict(tune.out$best.model,newdata=dat[-train,]))
```

```
> arrange(tune.out$performances,error)
```

	cost	gamma	error	dispersion
1	1e+00	0.5	0.07	0.08232726
2	1e+01	0.5	0.07	0.08232726
3	1e+00	1.0	0.07	0.08232726
4	1e+00	2.0	0.07	0.08232726
5	1e+00	3.0	0.07	0.08232726
6	1e+00	4.0	0.07	0.08232726
7	1e+01	3.0	0.08	0.07888106
8	1e+01	1.0	0.09	0.07378648
9	1e+01	4.0	0.09	0.07378648
10	1e+02	0.5	0.11	0.07278648

Test set performance

	pred	
true	1	2
1	67	10
2	2	21

**Test set error of best
model: 12%**

SVM Implementation

ROC Curves

```
library(ROCR)
```

```
rocplot=function(pred, truth, ...){  
  predob = prediction(pred, truth)  
  perf = performance(predob, "tpr", "fpr")  
  plot(perf,...)}
```

```
svmfit.opt=svm(y~., data=dat[train,], kernel="radial",gamma=0.5, cost=1,decision.values=T)  
fitted=attributes(predict(svmfit.opt,dat[train,],decision.values=TRUE))$decision.values  
rocplot(-fitted,dat[train,"y"],main="Training Data")
```

pred = vector of fitted values (e.g. -1.23, 1.05)

truth = vector of true labels (e.g. 1, 2)

decision.values=TRUE

**obtains fitted values from SVM
(e.g. -1.23, 1.05), see ISLR**

decision.values=FALSE

**obtains fitted class from SVM
(1 or 2)**


SVM Implem

ROC Curves

```
library(ROCR)
```

```
rocplot=function(predob = prediction(perf = performance(plot(perf,...))}
```

```
svmfit.opt=svm(y~.,  
fitted=attributes(pr  
rocplot(-fitted,dat
```



lues (e.g. -1.23, 1.05)
els (e.g. 1, 2)

st=1,decision.values=T)
UE))\$decision.values

=TRUE
lues