

# Bagging

**Algorithm class:** Non-parametric


**Mechanism:** Average predictions of many trees

**Applicable:** Both classification and regression problem


By model averaging, Bagging helps reduce variance and minimize overfitting

# The problem with single trees

Single shallow trees are poor predictors



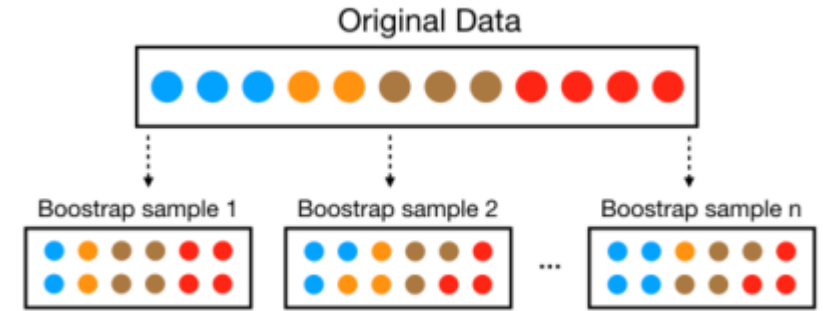
Single deep trees have high variance



Bagging uses this high variance to our advantage

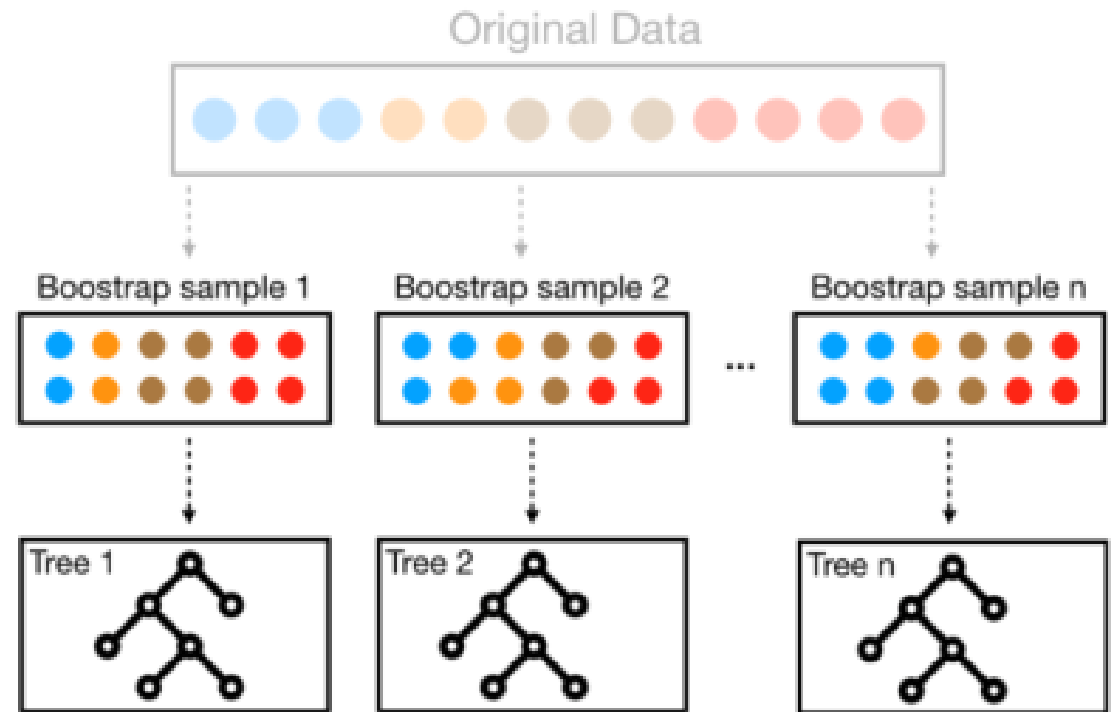
# Bootstrap Aggregating – wisdom of the crowd

1. Create bootstrap samples from training data
- 2.
- 3.




# Bootstrap Aggregating – wisdom of the crowd

1. Create bootstrap samples from training data
2. Fit an overgrown tree to each resampled dataset
- 3.




# Bootstrap Aggregating – wisdom of the crowd

1. Create bootstrap samples from training data
2. Fit an overgrown tree to each resampled dataset
3. Average predictions




# However, there is a problem


Decision Tree 1




Decision Tree 2




Decision Tree 3




Decision Tree 4



Decision Tree 5



Decision Tree 6



which prevents bagging from optimally reducing variance of the predictive values

# Random Forest

**Algorithm class:** Non-parametric

**Mechanism:** Average predictions of many trees (de-correlated)

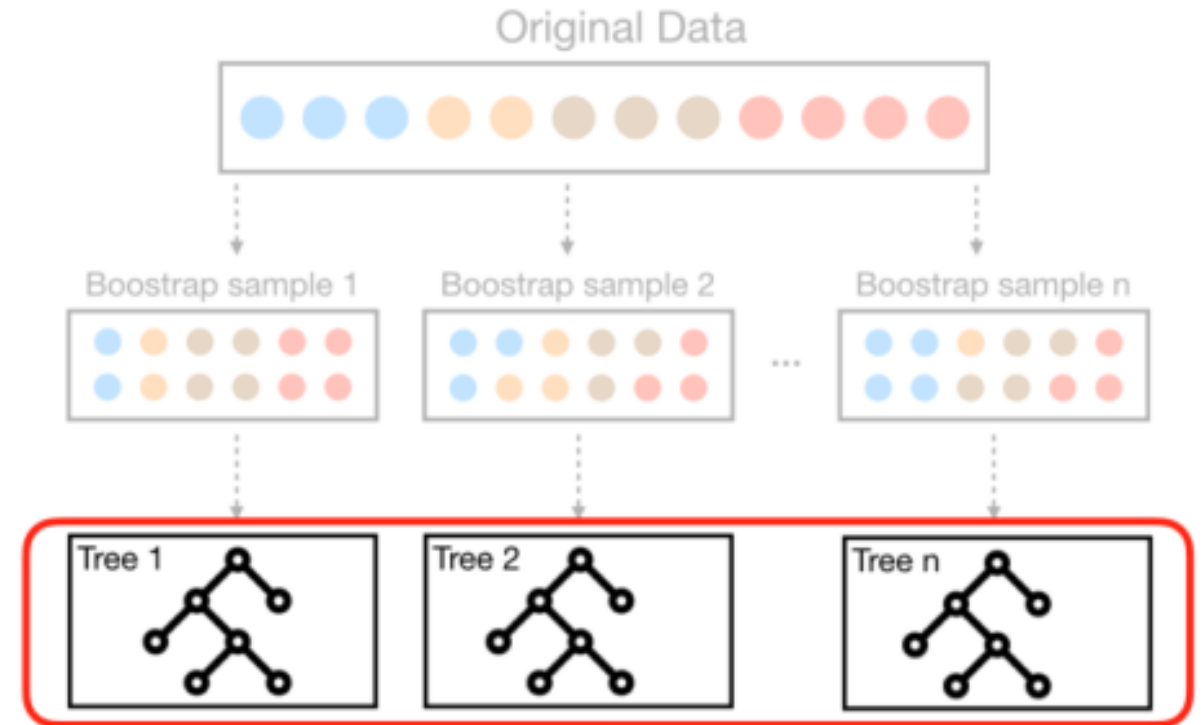
**Applicable:** Both classification and regression problem

Random Forest: generalization of Bagging, performance typically better

# Idea

Split variable randomization

- Follow a similar bagging process but ...



**Trees produced by bagging**

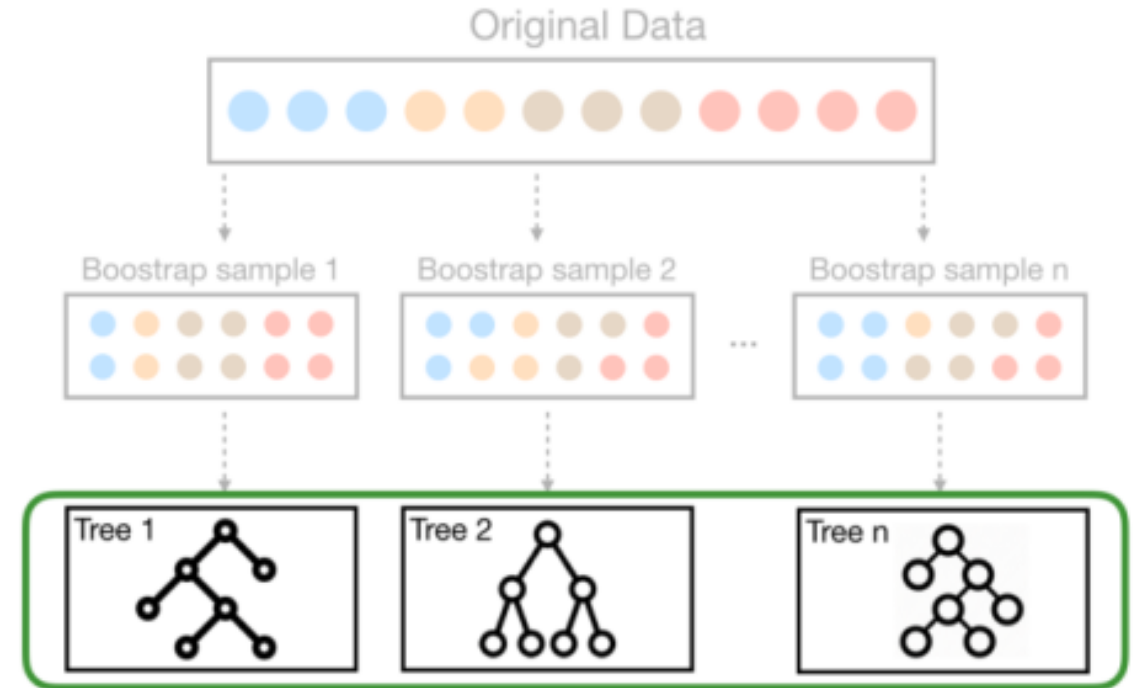


# Idea

## Split variable randomization

- Follow a similar bagging process but ...
- Each time **a split is to be performed**,

- regression trees:  $m = p/3$
- classification trees:  $m = \sqrt{p}$
- $m$  is commonly referred to as  $m_{try}$



**Trees produced by RF**

# Random Forest

Essentially

- Bagging introduces randomness into **rows of the data**
- Random forest introduces randomness into \_\_\_\_\_
  - This provides a more diverse set of trees that almost always lowers the prediction error

# Out of bag (OOB) Performance

- For large enough  $N$ , an original data point has a 63% probability of ending up in any bootstrap sample
  - i.e. 37% of the observations NOT used in for a particular tree
- These OOB observations: model performance assessment
- Let's now look at how to implement RF

# Implementation of Random Forest

- Simple way: ranger, full grid search
- More advanced: h2o, random grid search & early stopping rules

# Ames Housing Example (RF), with ranger package

Direct implementation of RF, no tuning

```
# number of features
n_features <- length(setdiff(names(ames_train), "Sale_Price"))

# train a default random forest model
ames_rf1 <- ranger(
  Sale_Price ~ .,
  ... data = ames_train,
  mtry = floor(n_features / 3),
  respect.unordered.factors = "order",
  seed = 123
)

# get OOB RMSE
(default_rmse <- sqrt(ames_rf1$prediction.error))
# 25488.39
```

For regression tree

Baseline RF model, RMSE  $\approx$  25,500

Next, we will look at how to tune hyperpara. to improve the model

# Random Forest

## Tuning Hyperparameters


Random forests provide good "out-of-the-box" performance but there are a few hyperparameters we can tune to increase performance.

# Trees


Mtry

Min node size/Max depth  
(Tree Complexity)

Sampling scheme



Typically have the largest impact on predictive accuracy



Some impact on predictive accuracy, but can increase computational efficiency

# Random Forest

## Tuning Hyperparameters: # Trees

Need to be sufficiently large: stabilize error rate

**Rule of thumb:** start with 10p trees and adjust as necessary

More trees provide robust and stable error estimates and variable importance measures

Computation time increases linearly with the number of trees

# Random Forest

Tuning Hyperparameters: mtry (#split vars)

Balance low tree correlation and reasonable predictive strength

**Rule of thumb default:**

Regression default:

Classification default:

Start with 5 values evenly spaced from 2 to p, including the default rule-of-thumb value

**Few relevant predictors: Should we  $\uparrow$  or  $\downarrow$  mtry?**



# Random Forest

Tuning Hyperparameters: Min node size/Max depth (Tree Complexity)

Control the complexity of individual trees

## **Rule of thumb:**

Regression default: 5

Classification default: 1

Start with 3 values (1,5,10)

If run time is a concern, can ↓run time substantially by ↑node size (tradeoff between runtime and model accuracy)

# Random Forest

## Tuning Hyperparameters: Sampling scheme

1. Sample size (default: 100%)
2. Sample with replacement / without replacement (default: with replacement)

### **Rationale:**

↓ Sample size reduces between-tree correlation  
(Smaller sample size produces more random trees)

### **Rule of thumb:**

3-4 values of sample sizes ranging from 25-100%  
Try both sampling with/without replacement

# Ames Housing Example (RF), with ranger package (cont'd)

## Tuning Strategy Illustration

```
# The code below searches across 120 combinations of hyperparameter settings.  
  
# create hyperparameter grid  
hyper_grid <- expand.grid(  
  mtry = floor(n_features * c(.05, .15, .25, .333, .4)),  
  min.node.size = c(1, 3, 5, 10),  
  replace = c(TRUE, FALSE),  
  sample.fraction = c(.5, .63, .8),  
  rmse = NA  
)
```

**Note:** `expand.grid` returns a dataframe with columns `mtry`, `min.node.size`, `replace`, `sample.fraction`, `rmse` (values to be filled)

# Ames Housing Example (RF), with ranger package (cont'd)

## Tuning Strategy Illustration

```
# execute full cartesian grid search
for(i in seq_len(nrow(hyper_grid))) {
  # fit model for ith hyperparameter combination
  fit <- ranger(
    formula      = Sale_Price ~ .,
    data         = ames_train,
    num.trees    = n_features * 10,
    mtry         = hyper_grid$mtry[i],
    min.node.size = hyper_grid$min.node.size[i],
    replace      = hyper_grid$replace[i],
    sample.fraction = hyper_grid$sample.fraction[i],
    verbose      = FALSE,
    seed         = 123,
    respect.unordered.factors = 'order',
  )
  # export OOB error
  hyper_grid$rmse[i] <- sqrt(fit$prediction.error)
}

# assess top 10 models
hyper_grid %>%
  arrange(rmse) %>%
  mutate(perc_gain = (default_rmse - rmse) / default_rmse * 100) %>%
  head(10)
```

**#trees**  
**mtry**  
**Node size**  
**Sample scheme**

**Fills rmse in hyper\_grid  
(created by expand.grid)**

# Ames Housing Example (RF), with ranger package (cont'd)

## Tuning Strategy Illustration

%improvement of RMSE w.r.t. baseline model

RMSE slightly improvement over  
baseline model

Observations

1. Default mtry = 26 (#features/3)  
nearly sufficient

2. Smaller node size performs better  
(deeper tree)

3. Sample <100% and sample  
without replacement consistently  
performs better

- Probably due to data having a lot  
of high-cardinality & imbalanced  
categorical features

	mtry	min.node.size	replace	sample.fraction	rmse	perc_gain
	1	26	1	FALSE	0.8 24713.06	3.041873
	2	26	3	FALSE	0.8 24847.98	2.512570
	3	20	3	FALSE	0.8 24917.05	2.241554
	4	20	1	FALSE	0.8 24929.10	2.194284
	5	32	5	FALSE	0.8 24940.14	2.150967
	6	32	1	FALSE	0.8 24978.78	1.999392
	7	32	3	FALSE	0.8 24990.83	1.952085
	8	26	5	FALSE	0.8 25004.10	1.900044
	9	20	5	FALSE	0.8 25028.46	1.804464
	10	12	1	FALSE	0.8 25029.93	1.798693

# Ames Housing Example (RF), with h2o package

## Benefits of h2o package:

- Random grid search
  - Full Cartesian hyperpara. search can be computationally expensive
  - Randomly jump from one random para. combination to another
- Can specify early stopping rules
  - E.g. #models trained  $\geq$  threshold, certain runtime elapses

```
# convert training data to h2o object  
train_h2o <- as.h2o(ames_train)
```

```
# set the response column to Sale_Price  
response <- "Sale_Price"
```

```
# set the predictor names  
predictors <- setdiff(colnames(ames_train), response)
```

# Ames Housing Example (RF), with h2o package

## Baseline h2o RF

- Syntax and result very similar to the baseline ranger RF

```
# The following fits a default random forest model with h2o
# Runtime: 1 minute on i7 CPU
h2o_rf1 <- h2o.randomForest(
  x = predictors,
  y = response,
  training_frame = train_h2o,
  ntrees = n_features * 10,
  seed = 123
)
```

h2o\_rf1

```
#H2ORegressionMetrics: drf
*** Reported on training data. **
#    ** Metrics reported on Out-Of-Bag training samples **
#
#    MSE: 626755219
#RMSE: 25035.08
#MAE: 15238.9
#RMSLE: 0.1415424
#Mean Residual Deviance : 626755219
```

**Similar to baseline RF using ranger**

# Ames Housing Example (RF), with h2o package

## h2o RF with Random Grid Search + Early Stopping Rule (Optional)

```
# To execute a grid search in h2o we need our hyperparameter grid to be a list.  
# For example, the following code searches a larger grid space than before with  
# a total of 240 hyperparameter combinations.  
  
# We then create a random grid search strategy that will stop if none of the  
# last 10 models have managed to have a 0.1% improvement in MSE compared to  
# the best model before that.  
# If we continue to find improvements then we cut the grid search off after 300 seconds
```

**Recall in ranger,**  
**we build the hyperpara. grid using the following syntax**

```
hyper_grid <- expand.grid(  
  mtry = floor(n_features * c(.05, .15, .25, .333, .4)),  
  min.node.size = c(1, 3, 5, 10),  
  replace = c(TRUE, FALSE),  
  sample.fraction = c(.5, .63, .8),  
  rmse = NA  
)
```



# Ames Housing Example (RF), with h2o package

## h2o RF with Random Grid Search + Early Stopping Rule (Optional)

```
# To execute a grid search in h2o we need our hyperparameter grid to be a list.  
# For example, the following code searches a larger grid space than before with  
# a total of 240 hyperparameter combinations.  
  
# We then create a random grid search strategy that will stop if none of the  
# last 10 models have managed to have a 0.1% improvement in MSE compared to  
# the best model before that.  
# If we continue to find improvements then we cut the grid search off after 300 seconds  
  
# hyperparameter grid  
hyper_grid <- list(  
  mtries = floor(n_features * c(.05, .15, .25, .333, .4)),  
  min_rows = c(1, 3, 5, 10),  
  max_depth = c(10, 20, 30),  
  sample_rate = c(.55, .632, .70, .80)  
)  
  
# random grid search strategy  
search_criteria <- list(  
  strategy = "RandomDiscrete",  
  stopping_metric = "mse",  
  stopping_tolerance = 0.001, # stop if improvement is < 0.1%  
  stopping_rounds = 10, # over the last 10 models  
  max_runtime_secs = 60*5 # or stop search after 5 min.  
)
```

# Ames Housing Example (RF), with h2o package

## h2o RF with Random Grid Search + Early Stopping Rule (Optional)

```
# To execute a grid search in h2o we need our hyperparameter grid to be a list.  
# For example, the following code searches a larger grid space than before with  
# a total of 240 hyperparameter combinations.  
  
# We then create a random grid search strategy that will stop if none of the  
# last 10 models have managed to have a 0.1% improvement in MSE compared to  
# the best model before that.  
# If we continue to find improvements then we cut the grid search off after 300 seconds
```

```
# hyperparameter grid
```

**In h2o, we use a list**

```
hyper_grid <- list(  
  mtries = floor(n_features * c(.05, .15, .25, .333, .4)),  
  min_rows = c(1, 3, 5, 10),  
  max_depth = c(10, 20, 30),  
  sample_rate = c(.55, .632, .70, .80)  
)
```

**Min node size**

```
# random grid search strategy  
search_criteria <- list(  
  strategy = "RandomDiscrete",  
  stopping_metric = "mse",
```

```
  stopping_tolerance = 0.001, # stop if improvement is < 0.1%  
  stopping_rounds = 10,      # over the last 10 models  
  max_runtime_secs = 60*5    # or stop search after 5 min.  
)
```

**Random grid-search strategy: “RandomDiscrete”**

- Randomly jump from one hyperpara. combination to another

**Early stopping criteria for grid-search**

- Stop if the last 10 RF models do NOT improve RMSE by 0.1%
- Stop if run time > 5 mins

# Feature Interpretation

For RF: 2 approaches for variable importance

At this point, do not need to know the details, just know there are 2 measures

Impurity (Same as CART)

- Based on the average total reduction in MSE

Permutation (Applicable for All ML models, will talk about it in more details)

- Permute a feature to a random value, see how it affects MSE

# Feature Interpretation

E.g. using ranger

```
# re-run model with impurity-based variable importance # re-run model with permutation-based variable importance
rf_impurity <- ranger(                                rf_permutation <- ranger(
  formula = Sale_Price ~ .,                            formula = Sale_Price ~ .,
  data = ames_train,                                    data = ames_train,
  num.trees = 2000,                                     num.trees = 2000,
  mtry = 32,                                             mtry = 32,
  min.node.size = 1,                                    min.node.size = 1,
  sample.fraction = .80,                                sample.fraction = .80,
  replace = FALSE,                                       replace = FALSE,
  importance = "impurity",                               importance = "permutation",
  respect.unordered.factors = "order",                  respect.unordered.factors = "order",
  verbose = FALSE,                                       verbose = FALSE,
  seed = 123                                              seed = 123
)
```

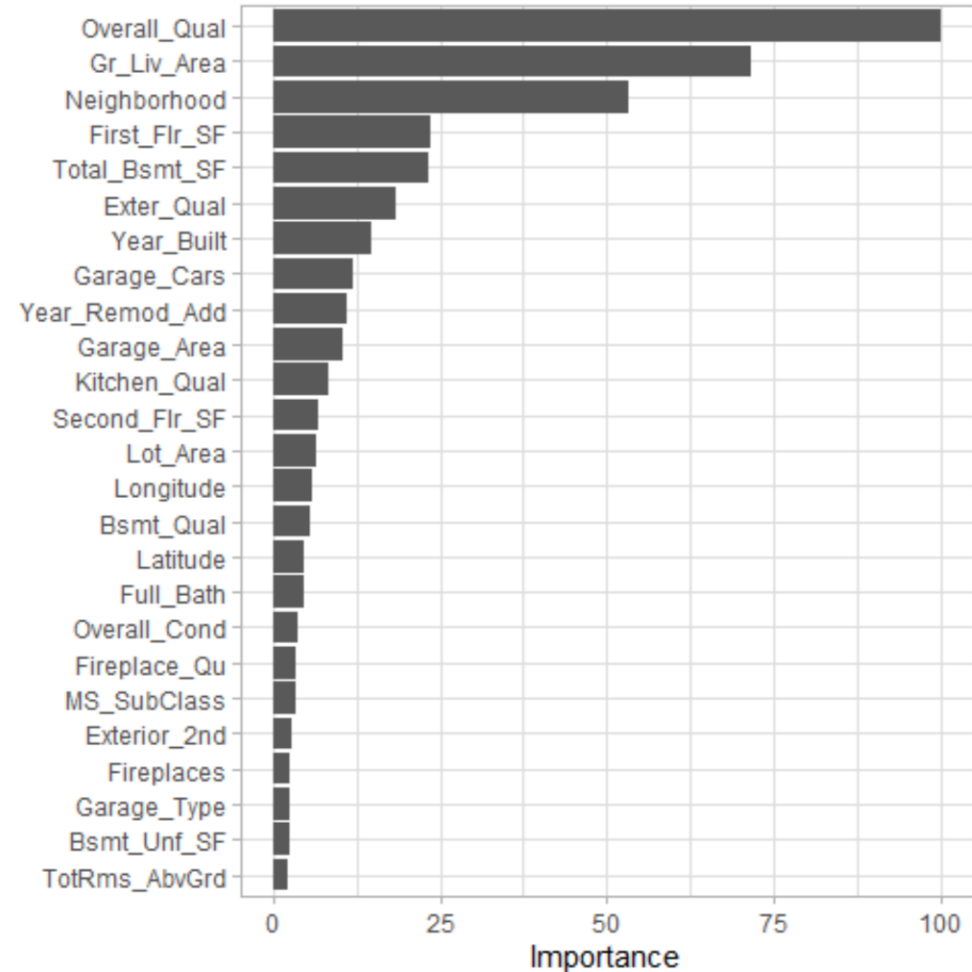

```
p1 <- vip::vip(rf_impurity, num_features = 25, scale = TRUE)
p2 <- vip::vip(rf_permutation, num_features = 25, scale = TRUE)
```

```
gridExtra::grid.arrange(p1, p2, nrow = 1)
```

# Feature Interpretation

Typically, similar variables at the top between the two approaches

- Can conclude top 3 important vars: Overall\_Qual, Gr\_Liv\_Area, Neighborhood



# Summary

Method	Hyperpara	Unique features	RMSE	Package Demonstrated	
CART	<ul style="list-style-type: none"> <li>Tree depth</li> <li>Node size</li> <li>cp</li> </ul>	Simple to interpret	-	rpart  caret method = "rpart"	
Random Forest	# Trees ( $\sim 10p$ ) Mtry (#split vars, $p/3$ or $\sqrt{p}$ ) Node size (Tree Complexity) Sampling scheme <ul style="list-style-type: none"> <li>Sample size</li> <li>Sample with/without replacement</li> </ul>	Subsample rows/cols Early Stopping (in adding trees)	$\sim 24000$	ranger  h2o Algorithm = "randomForest"	

End

# Ames Housing Example (RF), with h2o package

## h2o RF with Random Grid Search + Early Stopping Rule (Optional)

```
# perform grid search
# The following executes the grid search with early stopping turned on.
# The early stopping we specify below in h2o.grid() will stop growing an individual random
# experienced at least a 0.05% improvement in the overall OOB error in the last 10 trees.
# Runtime: 5 minutes
```

```
random_grid <- h2o.grid(
  algorithm = "randomForest",
  grid_id = "rf_random_grid",
  x = predictors,
  y = response,
  training_frame = train_h2o,
  hyper_params = hyper_grid,
  ntrees = n_features * 10,
  seed = 123,
  stopping_metric = "RMSE",
  stopping_rounds = 10,
  stopping_tolerance = 0.005,
  search_criteria = search_criteria
)
```

### Early stopping criteria for building one RF

- Stop if the last 10 trees added do NOT improve RMSE by 0.5%

```
# stop if last 10 trees added
# don't improve RMSE by 0.5%
# https://www.rdocumentation.org/packages/h2o/v
```



# Ames Housing Example (RF), with h2o package

h2o RF with Random Grid Search + Early Stopping Rule (Optional)

```
# collect the results and sort by our model performance metric of choice
random_grid_perf <- h2o.getGrid(
  grid_id = "rf_random_grid",
  sort_by = "mse",
  decreasing = FALSE
)
random_grid_perf
```

# Ames Housing Example (RF), with h2o package

h2o RF with Random Grid Search + Early Stopping Rule (Optional)

```
#H2O Grid Details
#=====
#
#   Grid ID: rf_random_grid
#Used hyper parameters:
#   - max_depth
#- min_rows
#- mtries
#- sample_rate
#Number of models: 66
#Number of failed models: 0
#
#Hyper-Parameter Search Summary: ordered by increasing mse
#max_depth min_rows mtries sample_rate model_ids mse
#1          30      1.0      20          0.8 rf_random_grid_model_57 6.0865378782043E8
#2          20      1.0      20          0.8 rf_random_grid_model_31 6.087217272346667E8
#3          20      1.0      26          0.8 rf_random_grid_model_32 6.179447344528593E8
#4          20      1.0      26          0.7 rf_random_grid_model_66 6.283513049703969E8
#5          30      1.0      32          0.632 rf_random_grid_model_64 6.373143927266866E8
#
#---
```

**Note: with early stopping, results may NOT be the same (#models searched in laptops of different speed will be different)**

**Assessed 66 models, best CV RMSE = 24670**

**This is near-optimal, and the random grid-search is more efficient**